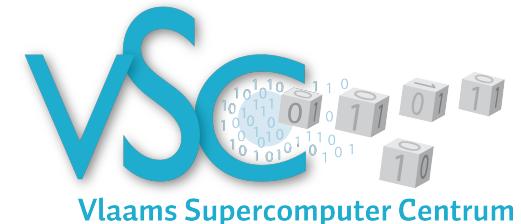


Code modernization a practical approach

HPC-TNT-1.2 fall 2016

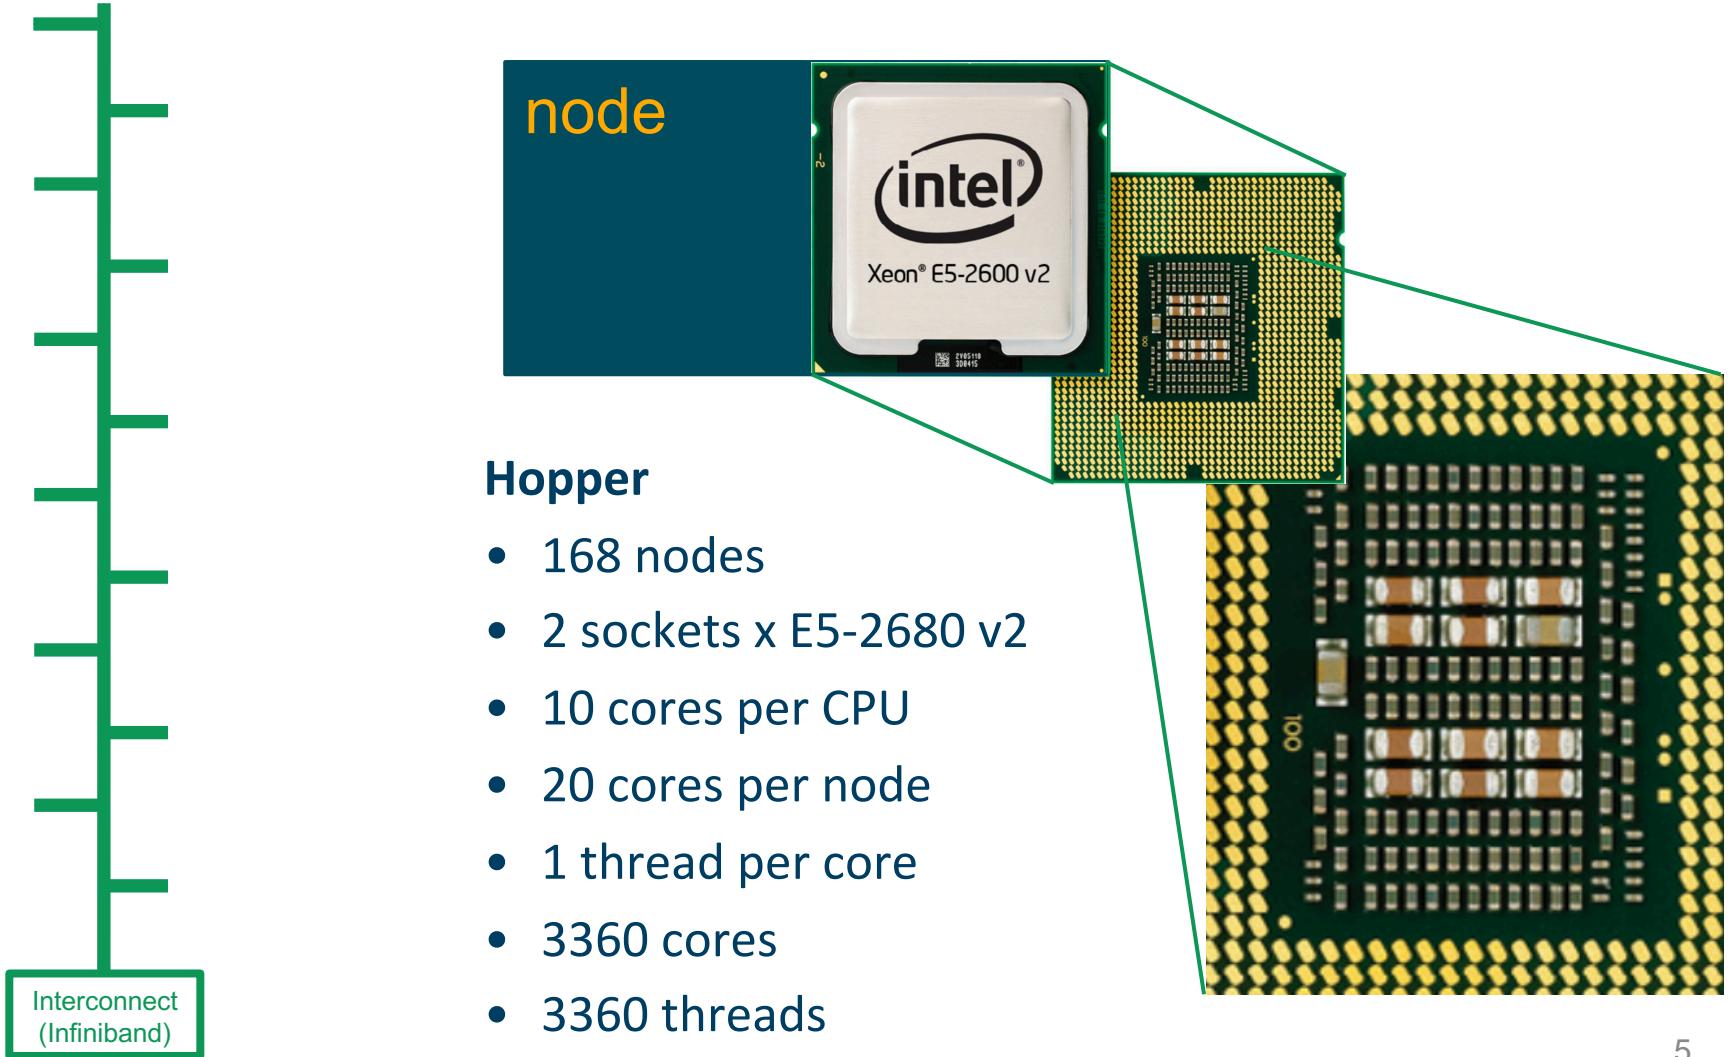


- A practical approach
 - To building efficient applications
 - To understanding performance issues
- To minimizing coding efforts
- To maximizing code flexibility

- Some computer architecture concepts related to performance and machine limits
 - Levels of parallelism → peak performance
 - Memory on modern CPUs → bandwidth, latency
- Example1 [toy problem]
 - Atomic system interacting through Lennard-Jones potential
 - Monte Carlo setting
- Example 2
 - Atomic system interacting through Lennard-Jones potential
 - Molecular dynamics setting
- Optimizing code \Leftrightarrow optimizing data access
 - Spatial sorting using space filling curve
- Choosing a programming language

A little terminology

- A (**compute**) **node** is basically a PC without all the peripheral devices
- Nodes are connected through a fast network: **interconnect**
- Every node has several **sockets**, each of which contains a processor
- Every processor contains many **cores**
- Each core can simultaneously execute one task: **thread**
 - Simultaneous multi-threading : more than one thread per core
 - Usually switched off on clusters
- Thread executes a sequential stream of instructions



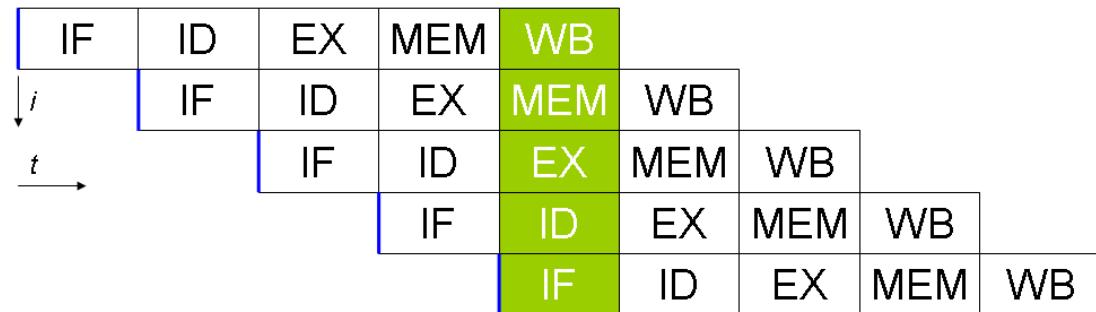
Levels of parallelism

- Several nodes can cooperate on a task
 - **Distributed** memory parallelization
 - Nodes communicate information via interconnect
 - Typically using MPI: Message Passing Interface
- Several cores can cooperate on a task
 - **Shared** memory parallelization
 - Nodes communicate information via memory
 - Often using OpenMP, but also MPI, ...

Levels of parallelism

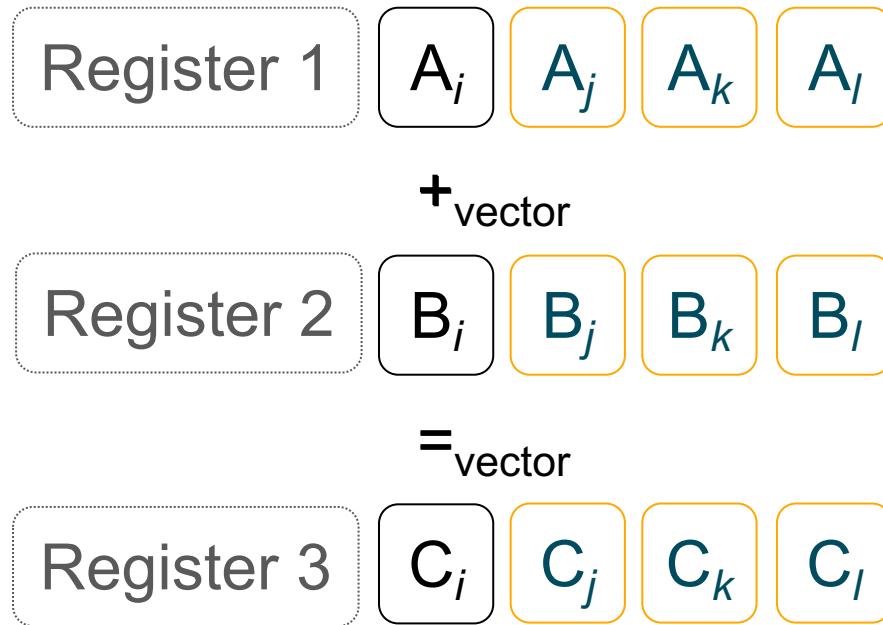
- A single core may exploit **pipelining** and **vectorisation** to execute instructions in parallel

- Instructions break down in micro-instructions
- Each micro-instruction uses a distinct part of the hardware
 1. Instruction fetch (IF)
 2. Instruction Decode (ID)
 3. Execution (EX)
 4. Memory Read/Write (MEM)
 5. Result Writeback (WB)



5 instructions executing simultaneously

Vectorization



Potentially 4x faster
[if the loads and stores can be executed fast enough]

Load 2×4 operands in register
 Execute 1 add instruction
 Store 1×4 results

Single Instruction
 Multiple Data (SIMD)

- Fused instructions: Fused Multiply-Add executes
 $y = a*x+b$
in one cycle
- Vector register width on Hopper is 256 bits
 - 8 single precision numbers
 - 4 double precision numbers

Levels of parallelism

- 3 levels of parallelism:
 - Intra-core: pipelining and SIMD
 -
 - Multi-core: shared memory
 - Multi-node: distributed memory
- Who does the work?
 - Compiler
 - (but it appreciates your help)
 - You and the compiler
 - OpenMP = directives
 - Relatively simple
 - You only
 - Harder

Case study – Lennard-Jones potential

1. Monte Carlo setting:

- compute the energy of configurations
- ensemble averages
- no forces
- no time integration

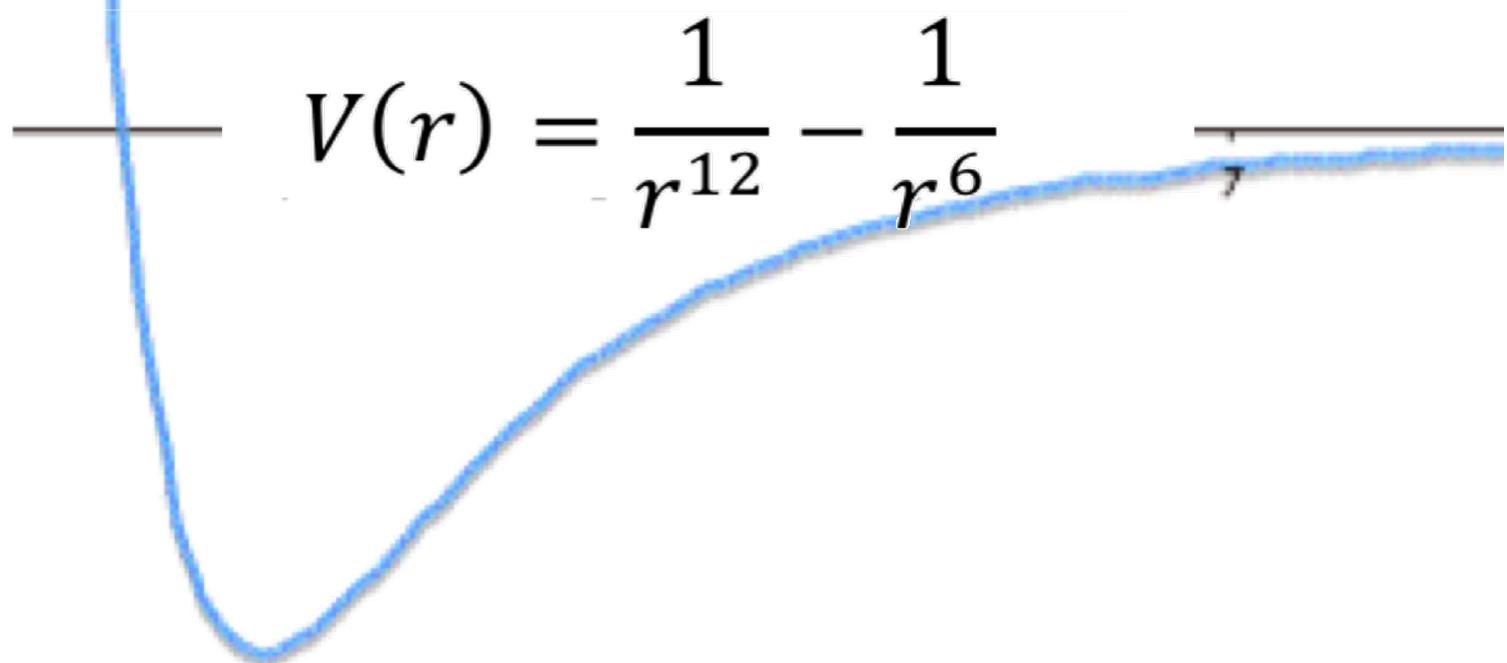
2. Molecular Dynamics setting:

- compute the time evolution of a collection of atoms
- forces
- [time integration]

- Compute interactions of an atom with **all** other atoms
- Doing this for all atoms is a $O(N^2)$ approach
 - but we do it only for a single atom just for the purpose of illustrating the behavior of the hardware
- Techniques to reduce to $O(N)$ will be discussed in MD setting
- We will only consider **single core** performance
 - That is the first thing to optimize anyway
 - Already complex enough for a single session
 - SIMD and pipelining are the only level of parallelism

Lennard-Jones interactions

- Lennard-Jones potential



Lennard-Jones potential

```
double VLJ0( double r ) {
    return 1./pow(r,12) - 1./pow(r,6);
} // 18.0 x slower
double VLJ1( double r )_{
    return std::pow(r,-12) - std::pow(r,-6);
} // 14.9 x slower
double VLJ2( double r ) {
    double tmp = std::pow(r,-6);
    return tmp*(tmp-1.0);
} // 7.8 x slower
double VLJ3( double r ) {
    double tmp = 1.0/(r*r*r*r*r*r);
    return tmp*(tmp-1.0);
} // 1.01 x slower
double VLJ( Real_t r ) {
    double rr = 1./r;
    rr *= rr;
    double rr6 = rr*rr*rr;
    return rr6*(rr6-1);
} // 1 x slower
```

Lennard-Jones potential

```
double VLJ( Real_t r ) {
    double rr = 1./r;
    rr *= rr;
    double rr6 = rr*rr*rr;
    return rr6*(rr6-1);
} // 1 x slower
double VLJ( Real_t r2 ) {
    double rr = 1./r2;
    // rr *= rr;
    double rr6 = rr*rr*rr;
    return rr6*(rr6-1);
} // avoid one sqrt per function call
```

Cost of instructions

- x_0, y_0, z_0 : coordinates of our central atom
- $x_1[1:m], y_1[1:m], z_1[1:m]$: coordinates of m neighbouring atoms
- Let $m=512, 1012, \dots, 2^{29} \sim 0.5 \cdot 10^9$
- surround by outer loop iterating $2^{29}/m$ times
 - every m -case executes 2^{29} evaluations of $V_{LJ}(r^2)$
- Variations
 - Loop over all m neighbouring atoms **contiguously**
 - Structure of arrays (SoA) : $x\ x\ x\ \dots\ y\ y\ y\ \dots\ z\ z\ z\ \dots$
 - Array of structures (AoS) : $x\ y\ z\ \dots\ x\ y\ z\ \dots\ x\ y\ z\ \dots$
 - Pick m random atoms in the arrays x_1, y_1, z_1

MC - Contiguous - SoA

```

integer :: m ! # of neighbour atoms
integer :: k ! # of iterations, m*k=cst, same amount of work per iteration
real(wp) :: x0,y0,z0, p(3*m)
! Contiguous access, SoA: p=[xxx...yyy...zzz...]
do ik=1,k
    do im=1,m
        r2 = (p(im)-x0)**2 +(p(m+im)-y0)**2 +(p(2*m+im)-z0)**2
        v = v + lj_pot2(r2)
    enddo
enddo

```

1	x					y					z				
2		x				y					z				
3			x			y					z				

```

integer :: m ! Number of neighbour atoms
real(wp) :: x0,y0,z0, p(3*m)
! ordered access, AoS: p=[xyzxyzxyz...]
do ik=1,k
  do im=1,m
    r2 = (p(im)-x0)**2 +(p(1+im)-y0)**2 +(p(2+im)-z0)**2
    v = v + lj_pot2(r2)
  enddo
enddo

```

MC - Random access

```

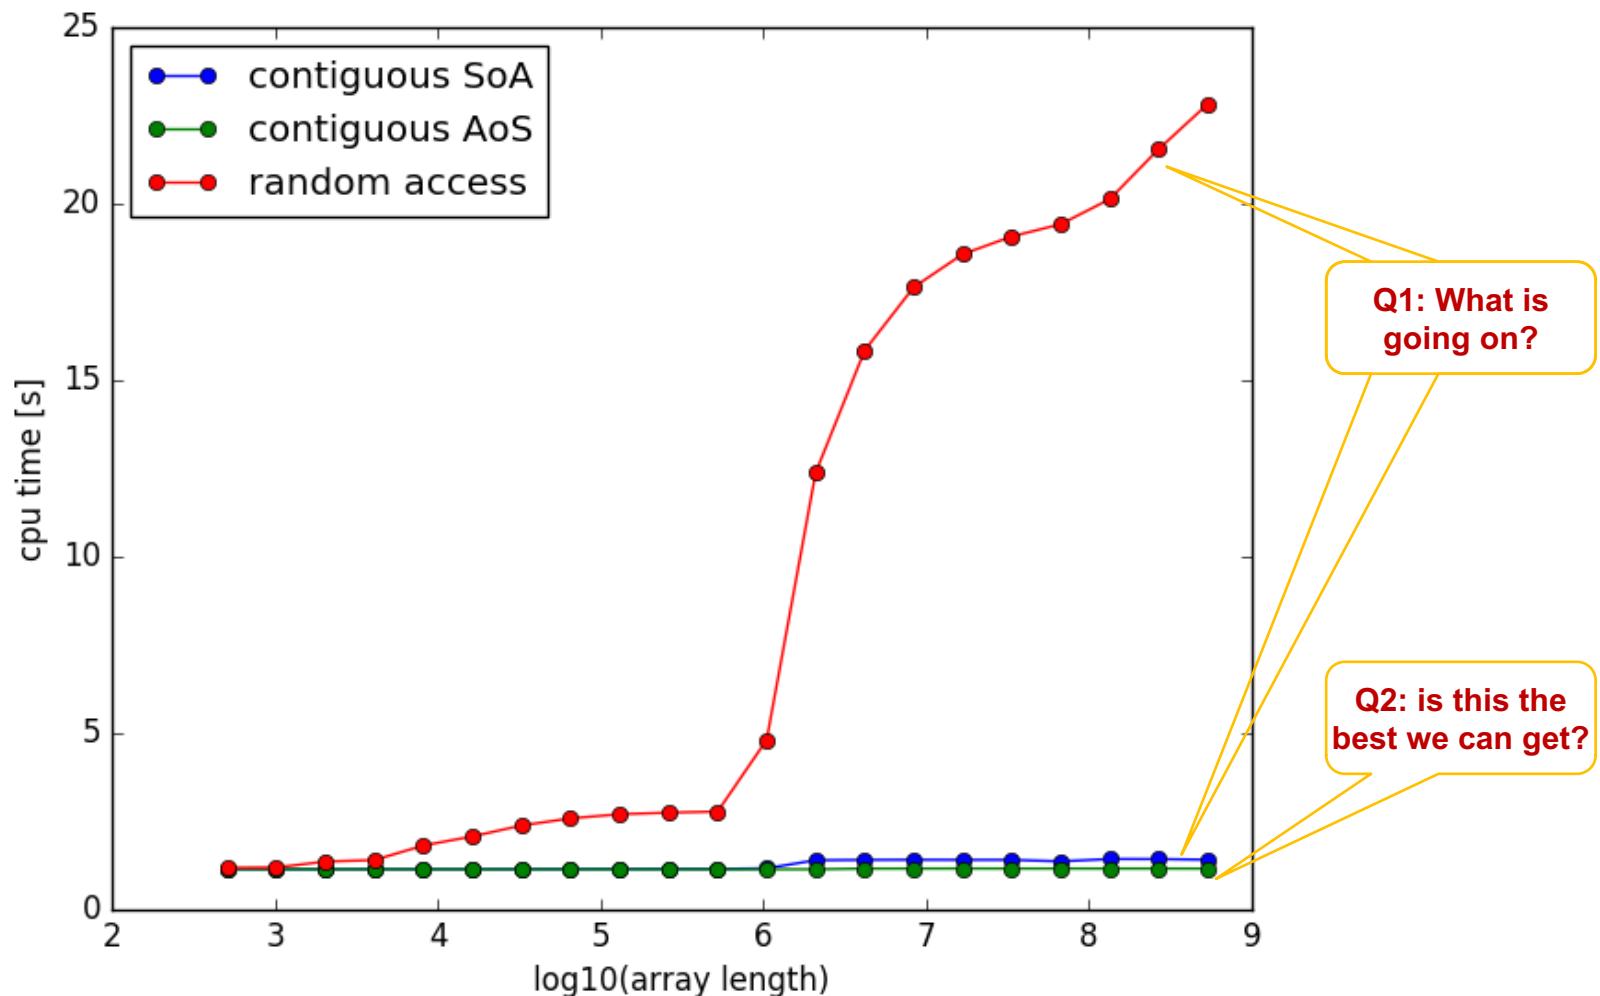
integer :: m ! Number of neighbour atoms
real(wp) :: x0,y0,z0, p(3*m)

integer :: j(m) ! random permutation of [1:m]

! random access
do ik=1,k
    do im=1,m
        r2 = (p(j(im))-x0)**2 +(p(j(im)+m)-y0)**2 +(p(j(im)+2*m)-z0)**2
        v = v + lj_pot2(r2)
    enddo
enddo

```

1			x				y					z		
2		x					y					z		
3				x				y						z



- Which factors influence performance of a code?
- Machine limits

- $\stackrel{\text{def}}{=}$ Maximum # floating point operations per second
- For a single core the peak performance =
 - 2*8 instructions per cycle in SP
 - 2*4 instructions per cycle in DP
 - The 2 comes from the fused multiply and add
 - The 8, resp. 4 come from the vector register width
- Peak performance per node
 - (1 cycle = 1/clock_frequency)
 - Assuming 1 hardware thread per core:
 - (#cores=20) * 2*8(SP) * (f=2.8Ghz) = 896 Gflop/s
 - (#cores=20) * 2*4(DP) * (f=2.8Ghz) = 448 Gflop/s

- Peak performance is not the only limiting factor...
- It is not the most common limiting factor
- Instructions operate on data, ...
- Data resides in memory
- Accessing data takes time (and energy)
 - Data has to be moved from memory to cpu register before it can be processed
- Peak performance has improved much faster than the speed at which data can be moved between memory and cpu

Speed of moving data

- Memory **bandwidth**

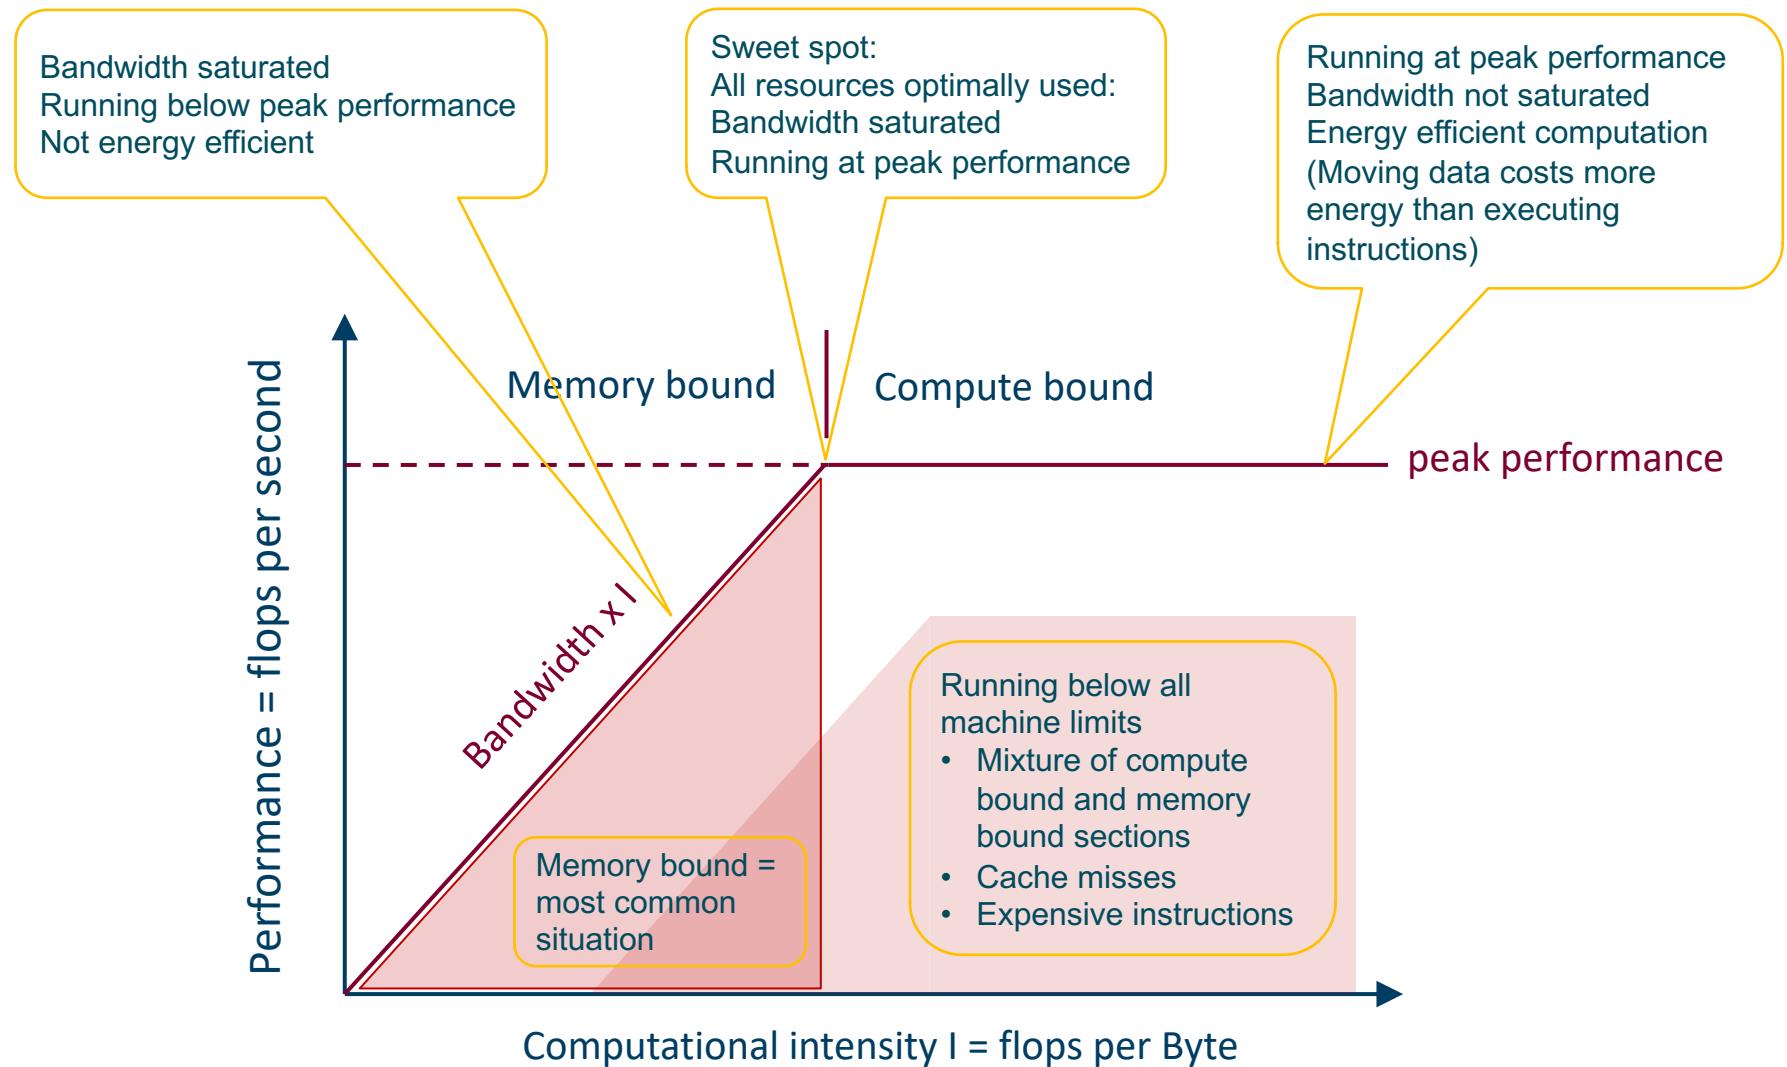
- The maximum number of bytes that can be moved per second between main memory and the cores
- Hopper
 - 92-110 MB/s (varies depending on read:write ratio)

- Memory **latency**

- The number of cycles (or the time) needed to fetch a single item from main memory
- Hopper
 - ~180 cycles (on socket)
 - ~350 cycles (across sockets)

- Code is **compute bound** if
 - The cpu can execute its compute instructions without having to wait for data
 - The limit is the theoretical peak performance
 - [Used to be the common case – not any more]
- Code is **memory bound** if
 - A considerable amount of cycles is spent waiting for data
 - Too much data requested:
 - **Bandwidth saturation** = machine limit
 - Too distant data requested:
 - If data is not in the cache: latency penalty
 - Latency problem = machine limit
 - [most common situation]

Roofline model



Memory bound - consequences

- Optimizing code **was** about organizing compute instructions
 - Pretty straightforward: less compute cycles is less cputime
 - Algorithmic complexity was important guideline
- Optimizing code is optimizing data access
 - To keep the processor busy doing useful stuff
 - **Algorithmic complexity is no longer a guarantee** for optimal performance
 - E.g. linear search (as in a map) often faster than binary or other search algorithms, also sorting
 - For large N low order complexity wins, but hardware caching takes an early lead
 - Understanding how memory works is necessary
 - Experimenting and measuring is necessary

Hierarchical memory organisation

On-chip (E5-2680 v2)

ALU

Registers: ~1kB per core 0 cycles

L1 Cache: 32 kB per core ~1 cycles

L2 Cache: 256 kB per core ~10 cycles

L3 Cache: 25 MB per socket ~50 cycles

Off-chip

DRAM: 64-256 GB per node ~200 cycles

memory
size speed



Memory organisation

- Memory is not fetched on a per item basis
- But in chunks called **cache lines**
 - typically 64 Bytes long
 - 16 single precision items
 - 8 double precision items

- **Linear search** of array $A[i]$, $i=1..n$
 - $A[1]$ is not in cache, wait time before cache line is loaded, dram latency (~ 200 cycles) and before item $A[1]$ can be examined
 - Once the cache line is loaded, $A[2..16]$ are also in L1 cache and are ready to be processed without delay
 - Hardware recognizes your loop over the array and keeps loading next (or previous) cache lines into the L1 cache, so that the delay is vanishing
 - Depending on how much work it takes to examine each item, as soon as item $A[16]$, the next cache line $A[17:32]$ may have been loaded already or not
 - In any case, the wait time is now less than the dram latency (~ 200 cycles)
 - The limitation becomes memory bandwidth of the machine

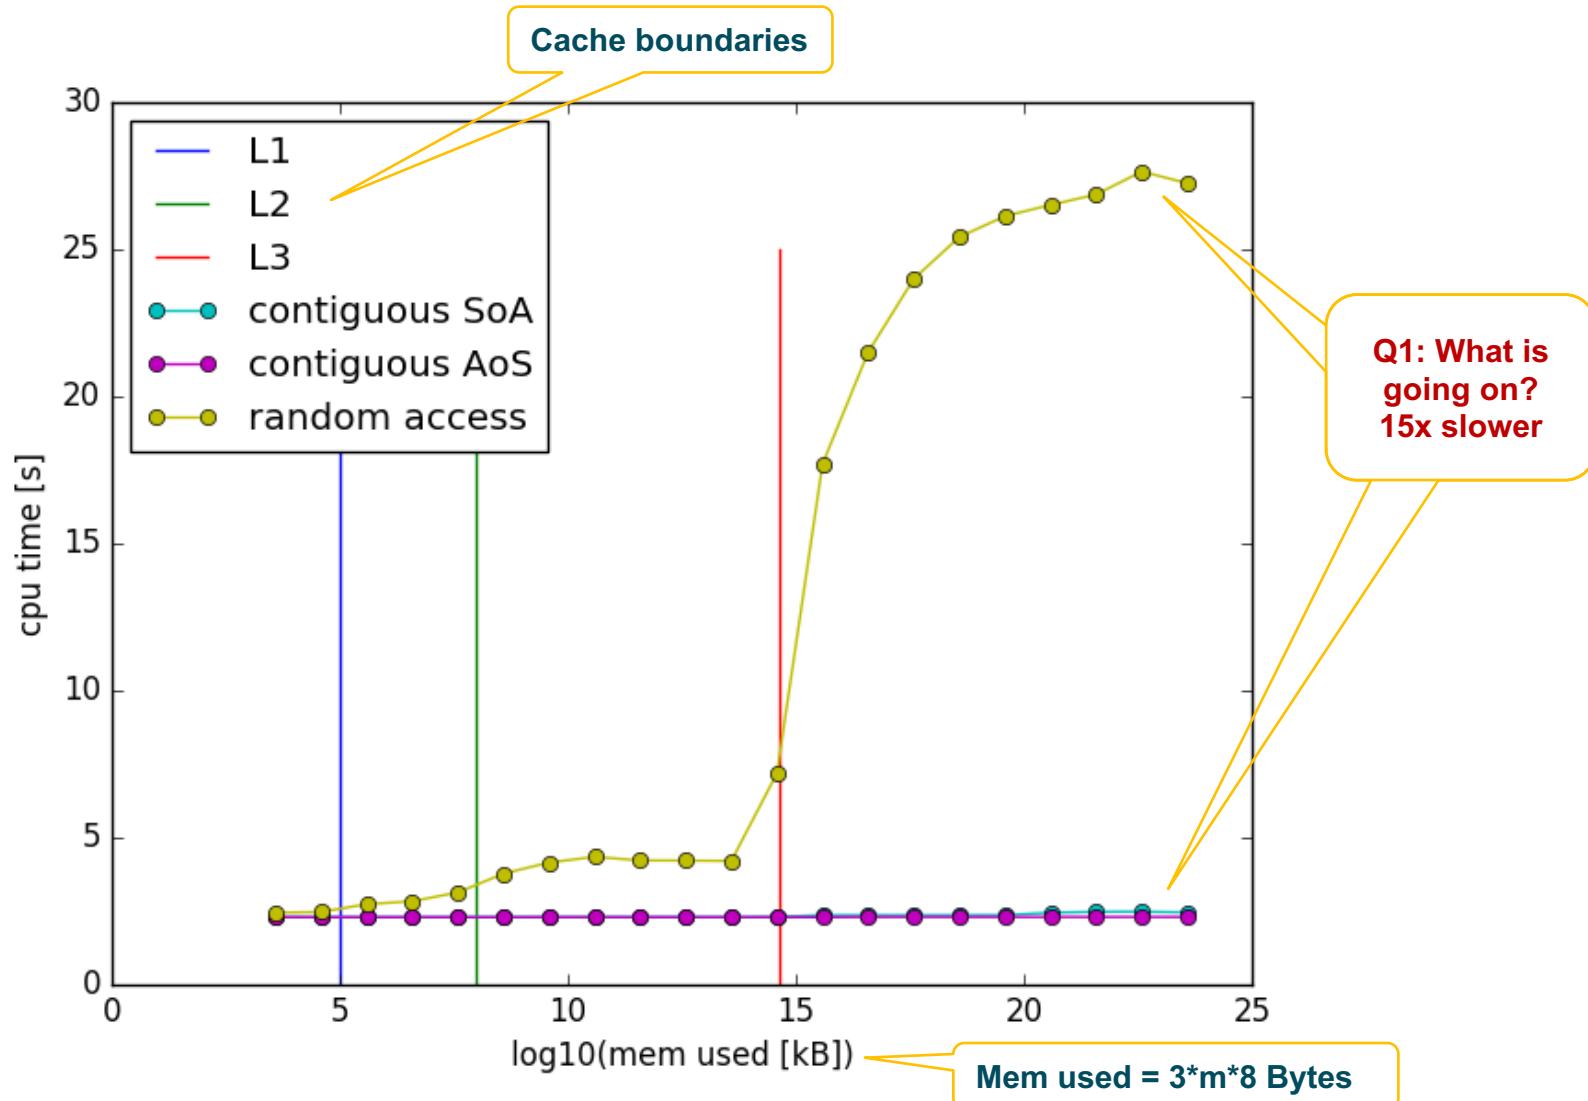
- **Binary search** of array $A[i]$, $i=1..n$
 - $A[n/2]$ is not in cache, wait time before cache line is loaded, dram latency
 - Next item needed is $A[n/4]$ or $A[3n/4]$, which is not in the cache, dram latency hits you again
 - In fact, the dram latency keeps on hitting you until the search range is reduced to one or two cache lines,
 - You do only one examination/dram latency, as opposed to 16/dram latency in linear search.

Memory organization

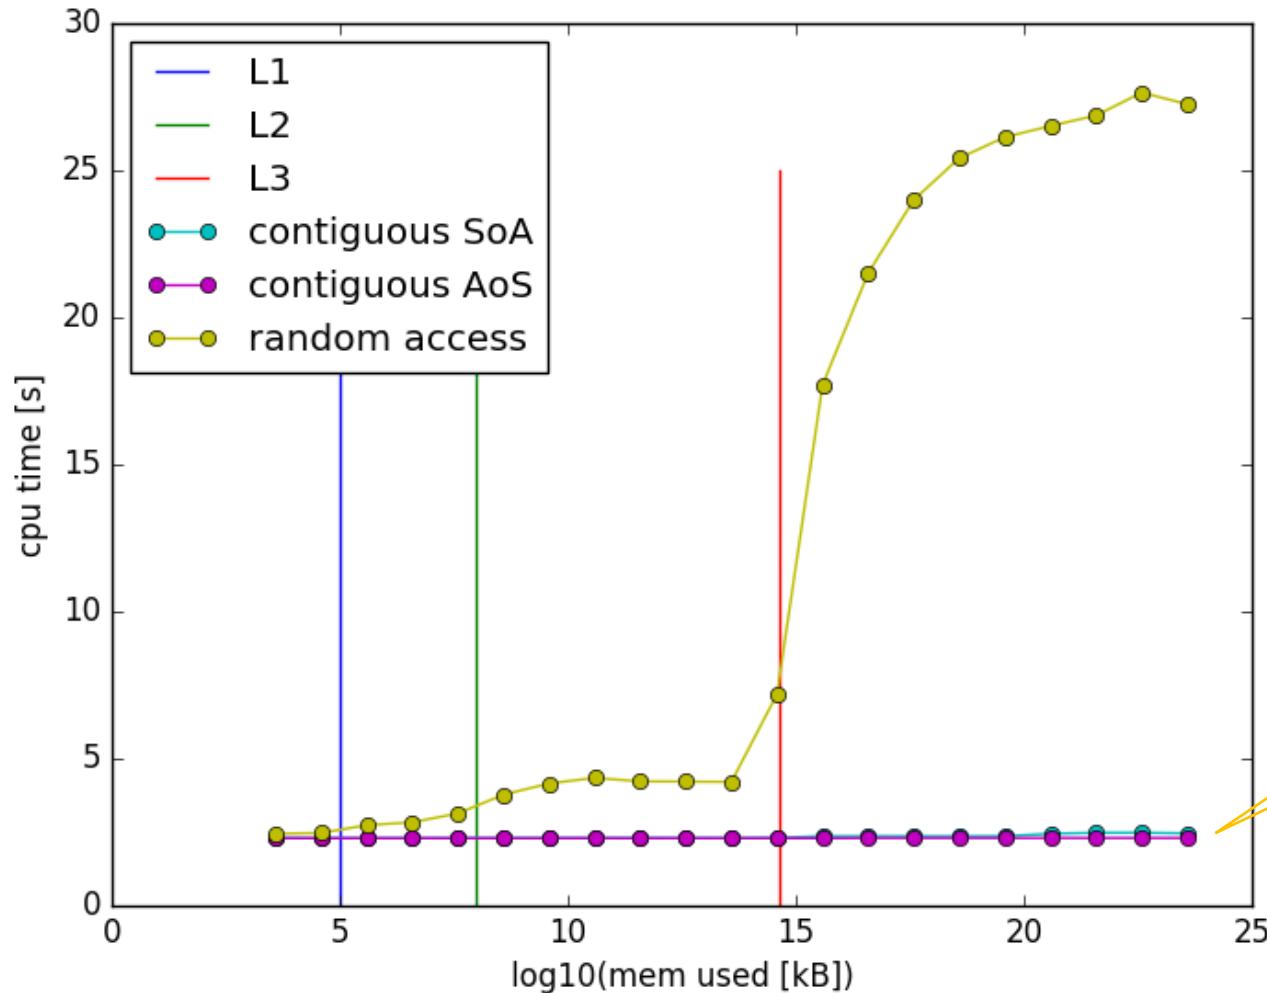
- Instructions are also data stored in memory
- Branching instructions can cause cache misses too!
 - Instruction cache misses
- Avoid unpredictable branches in loops

- code::dive conference 2014 - Scott Meyers: Cpu Caches and Why You Care
- <https://www.youtube.com/watch?v=WDIkqP4JbkE>

Experiment



Experiment

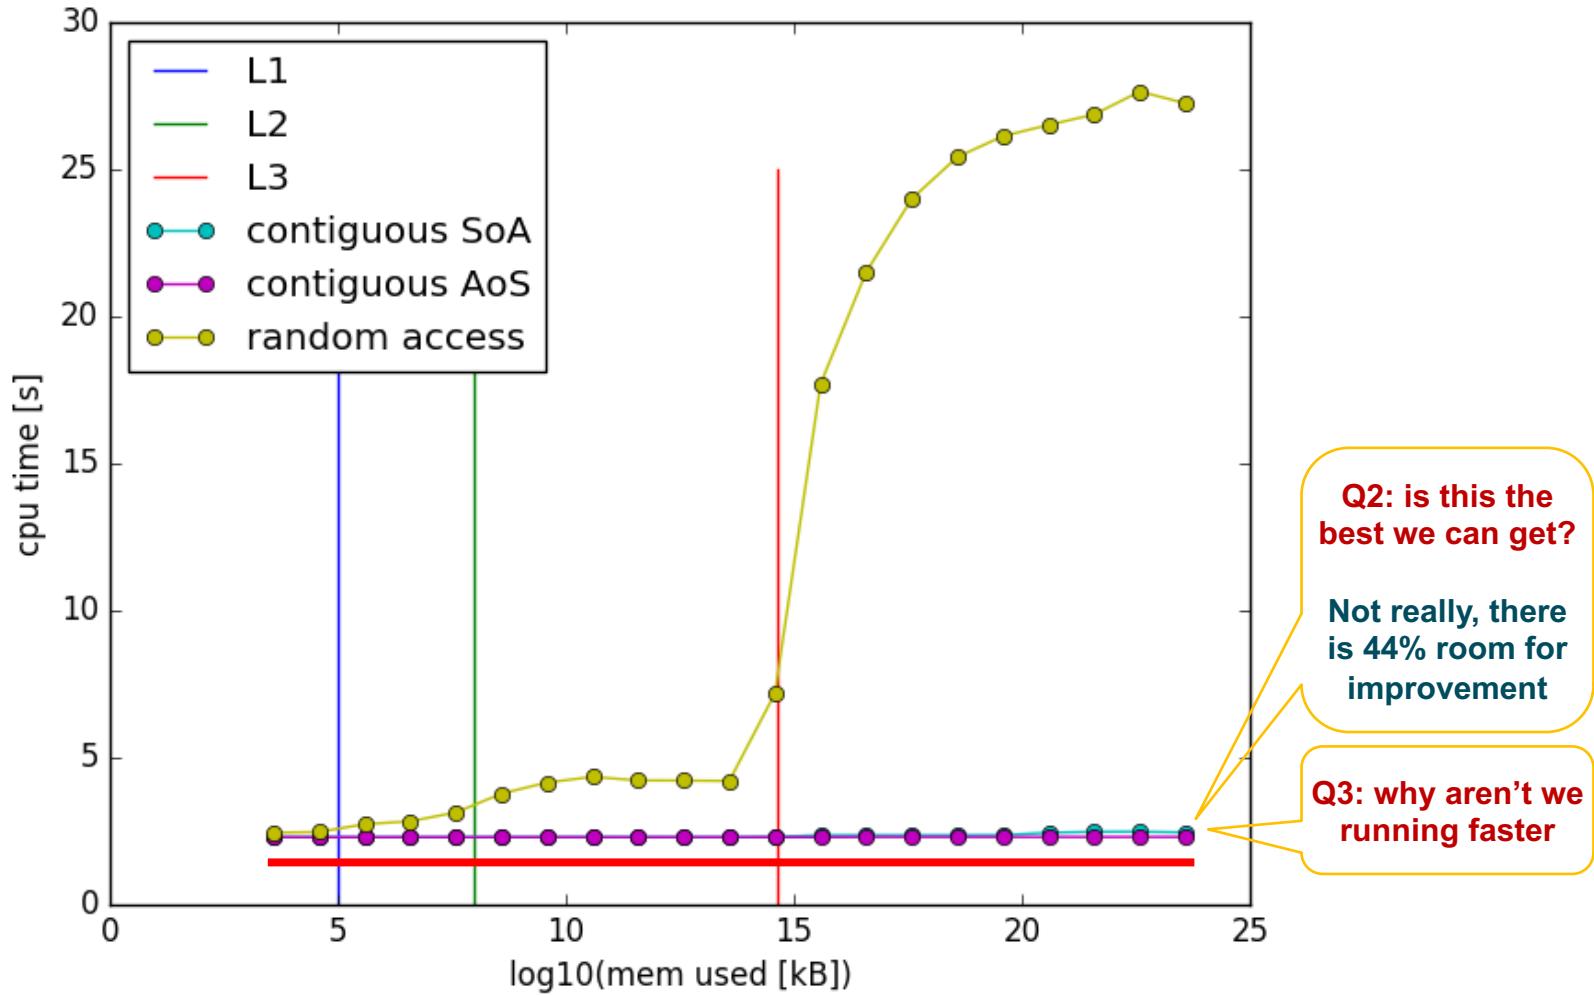


Q2: is this the best we can get?

Flops per second

- ! Contiguous access, SoA: $p=[xxx...yyy...zzz...]$
- do ik=1,k
 - do im=1,m
 - $r2 = (p(im)-x0)^{**2}$
 - $+(p(m+im)-y0)^{**2}$
 - $+(p(2*m+im)-z0)^{**2}$
 - 3-, 2+, 3*
 - ! $r = lj_pot2(r)$
 - $r2i = 1.0d0/r2$
 - $rr6i = r2i*r2i*r2i;$
 - $lj_pot2 = 4.0d0*rr6*(rr6-1.0d0);$
 - 1/
 - 2*
 - 2*, 1-
 - -----
 - 14 flops
 - enddo
- enddo
- 14 flops * 2^{29} iterations in 1.2 s = $6.26 \cdot 10^9$ flops/s
- peak performance:
 $1*1*4*2.8 \text{ GHz} = 11.2 \text{ Gcycles/s} = 11.2 \text{ Gflops/s}$
- We are running at 55.9 % of peak performance

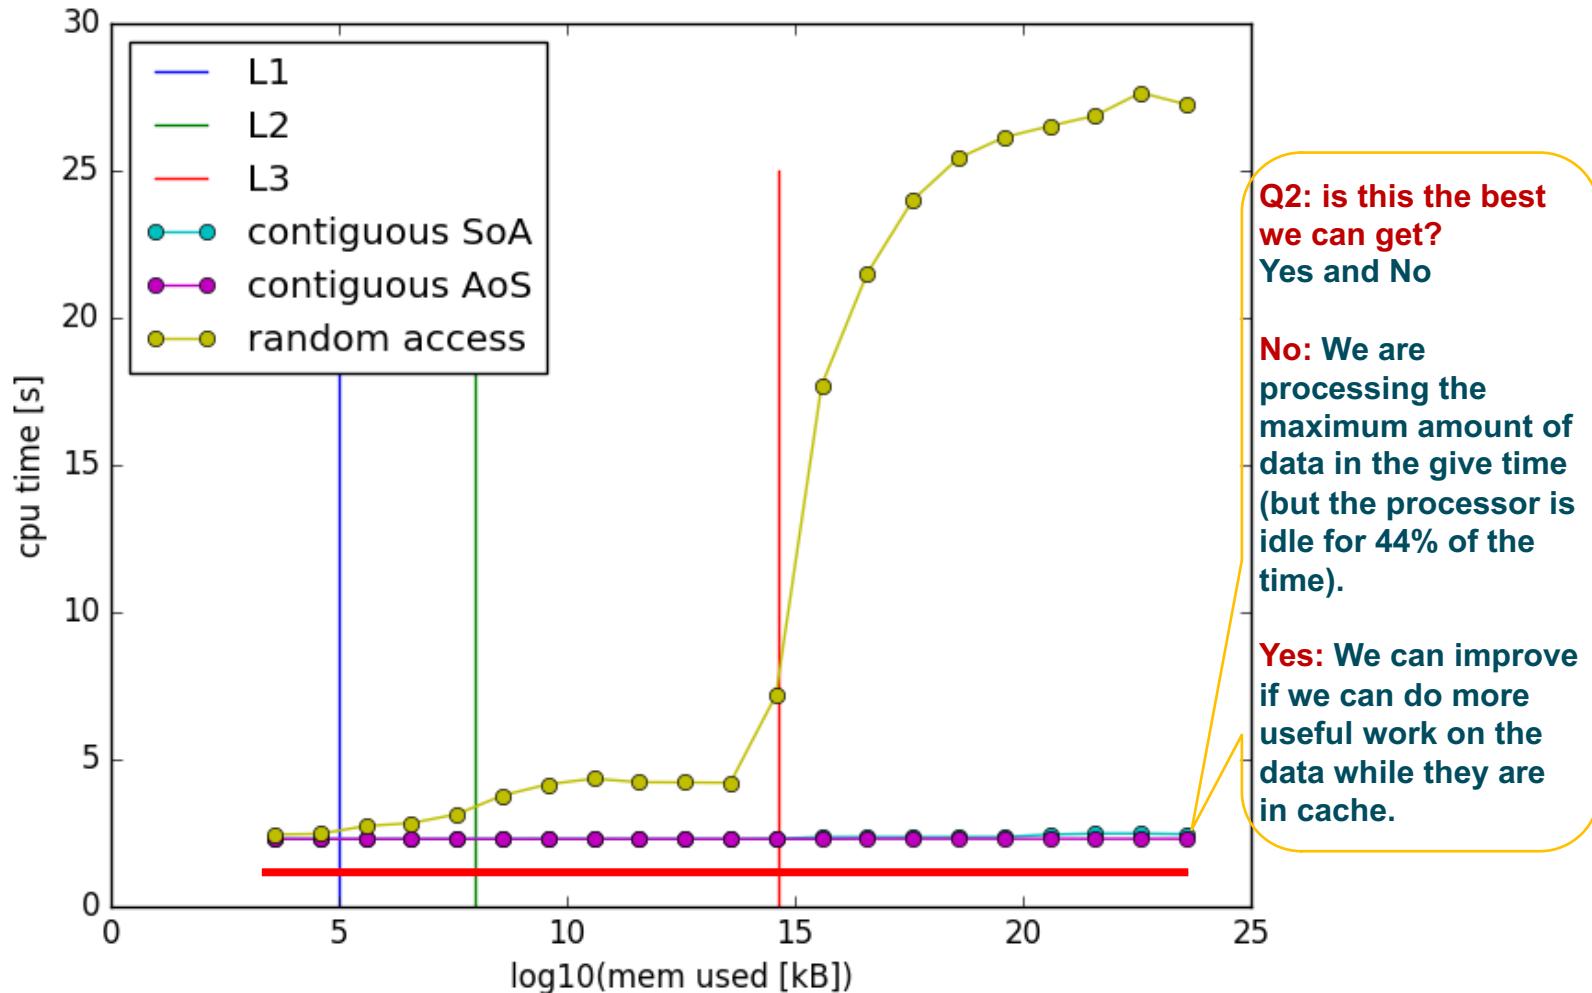
Experiment



Bytes per second

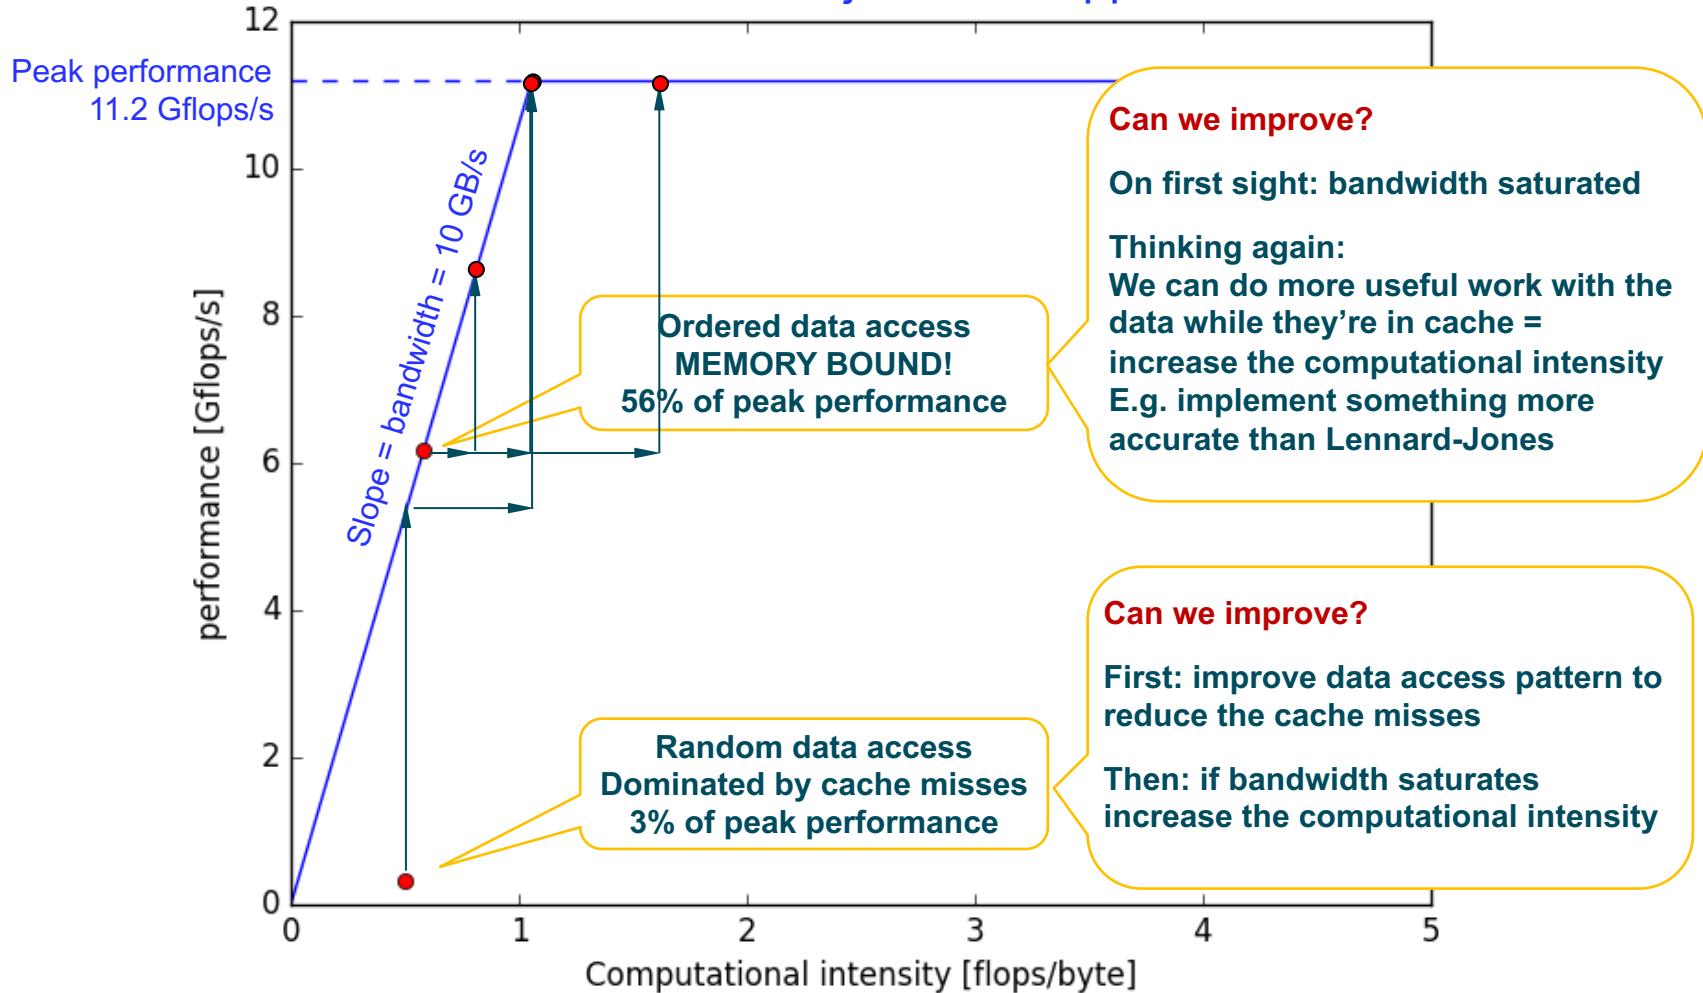
- ! Contiguous access, SoA: $p=[xxx...yyy...zzz...]$
- do ik=1,k
 - do im=1,m
 - $r2 = (p(im)-x0)^{**2}$
 - $+(p(m+im)-y0)^{**2}$
 - $+(p(2*m+im)-z0)^{**2}$
 - 3-, 2+, 3* • 3 DP
 - ! $r = lj_pot2(r)$
 - $r2i = 1.0d0/r2$
 - $rr6i = r2i*r2i*r2i;$
 - $lj_pot2 = 4.0d0*rr6*(rr6-1.0d0);$
 - enddo
- enddo
- $24 \text{ B} * 2^{29} \text{ iterations in } 1.2 \text{ s} = 10.7 \text{ GB/s}$
- Bandwidth measured by Intel mlc:
 - 109 GB/s for 10 threads (all reads)
 - 10.9 GB/s for 1 thread
- We are running at maximal bandwidth
 - Bandwidth saturation

Experiment



Roofline model

Roofline for a 1 core job on a hopper node



- For simple cases
 - a back of the envelope calculation like this
 - and an understanding of how memory works can guide you to more efficient code
- For real cases we need something more sophisticated

- Intel Advisor xe
 - Vectorization and threading
 - Intel VTune Analyzer xe
 - Data access and cpu utilization
 - Intel Inspector
 - Thread performance analysis (OpenMP, Intel TBB)
 - Intel Cluster Inspector
 - MPI process performance analysis
- For a later session**

- <https://software.intel.com/en-us/get-started-with-advisor>
- <https://software.intel.com/en-us/get-started-with-vtune>

```
$ ssh -X vsc20170@login.hpc.uantwerpen.be  
Last login: Thu Sep  8 16:38:25 2016 from 143.169.185.55
```

```
Welcome to Hopper!
```

```
...
```

```
vsc20170@ln02 ~$
```

```
vsc20170@ln02 ~$ qsub -I -X
```

Allow X11 forwarding. On mac os x install Xquartz, On windows install Xming

```
vsc20170@r5c6cn05 ~$
```

Start interactive job with X11 forwarding.

```
vsc20170@r5c6cn05 ~$ module load Advisor
```

```
vsc20170@r5c6cn05 ~$ module list
```

```
Currently Loaded Modulefiles:
```

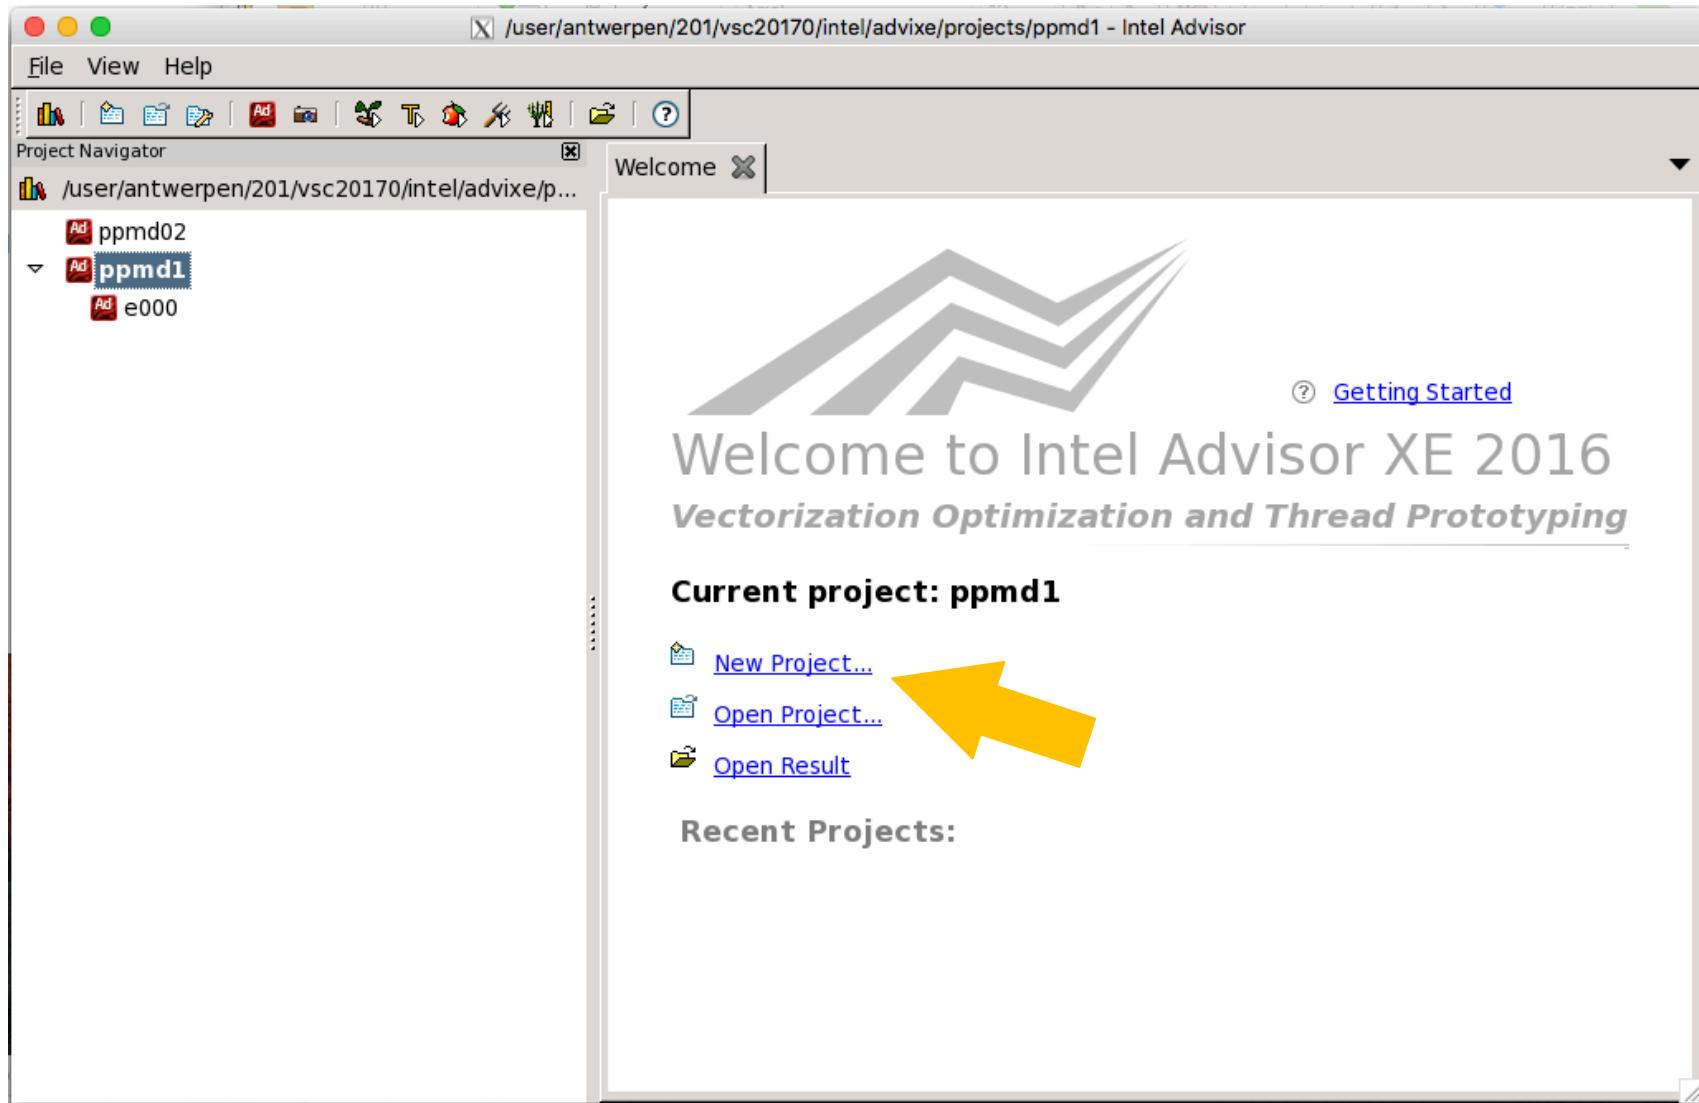
```
1) Advisor/2016_update4
```

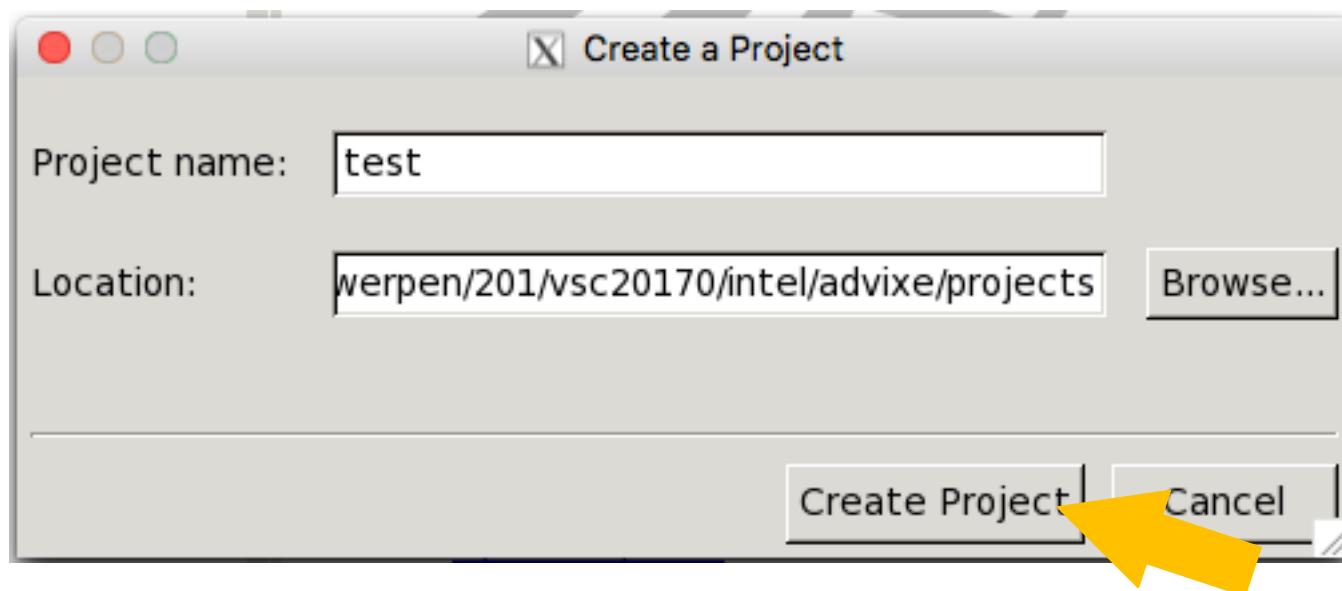
```
vsc20170@r5c6cn05 ~$
```

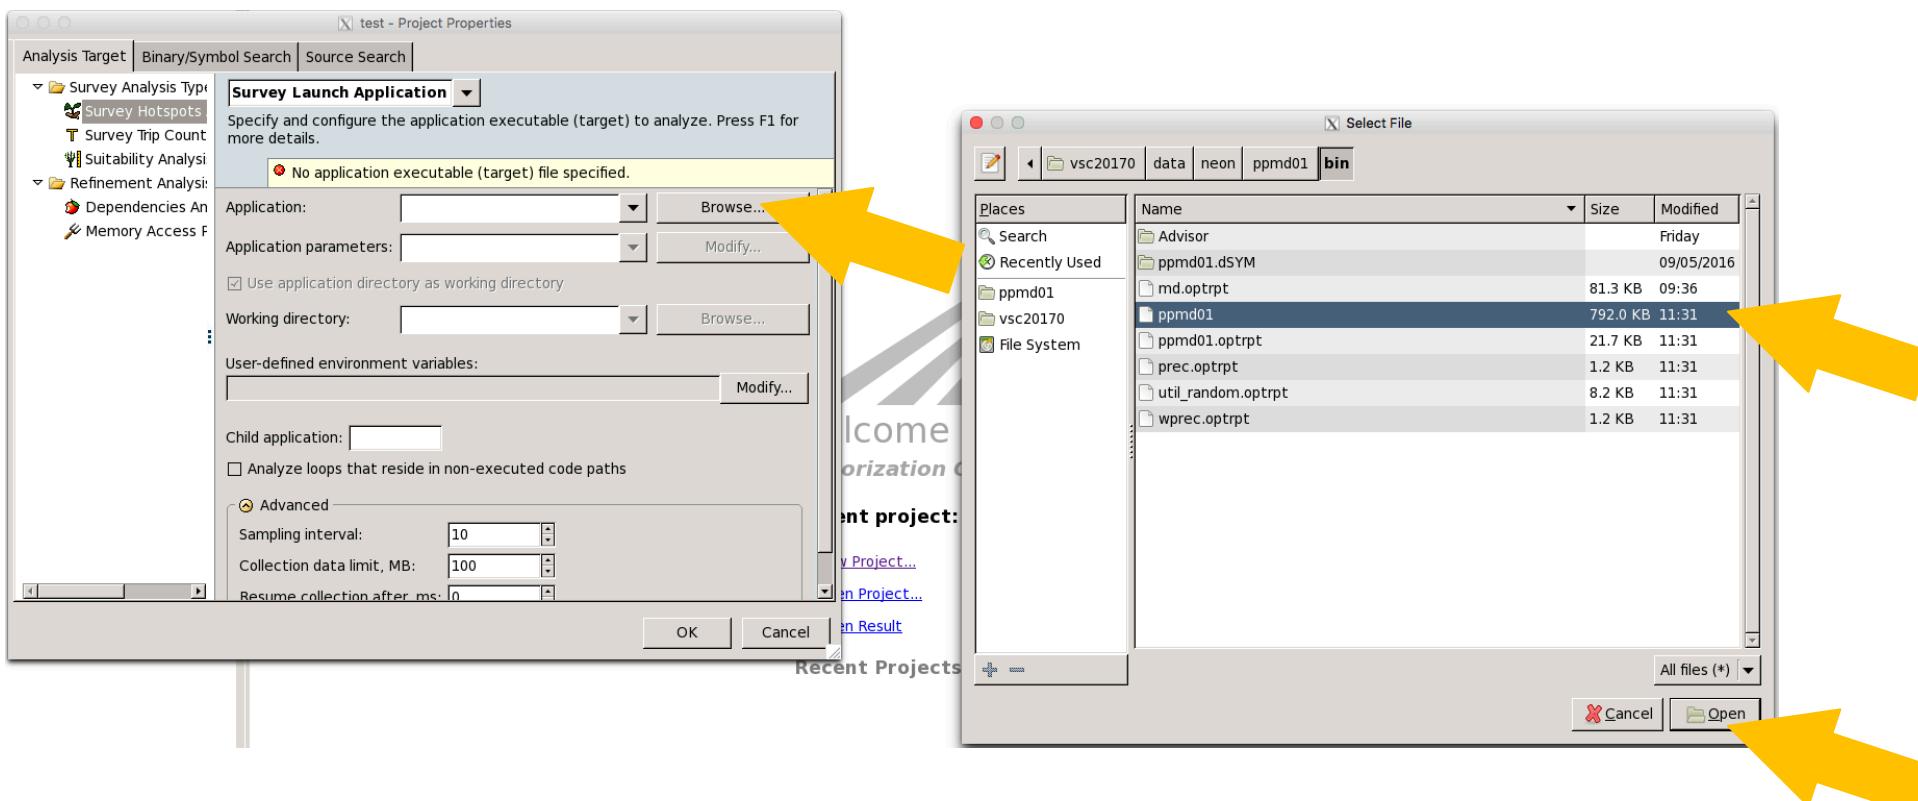
```
vsc20170@r5c6cn05 ~$ advixe-gui &
```

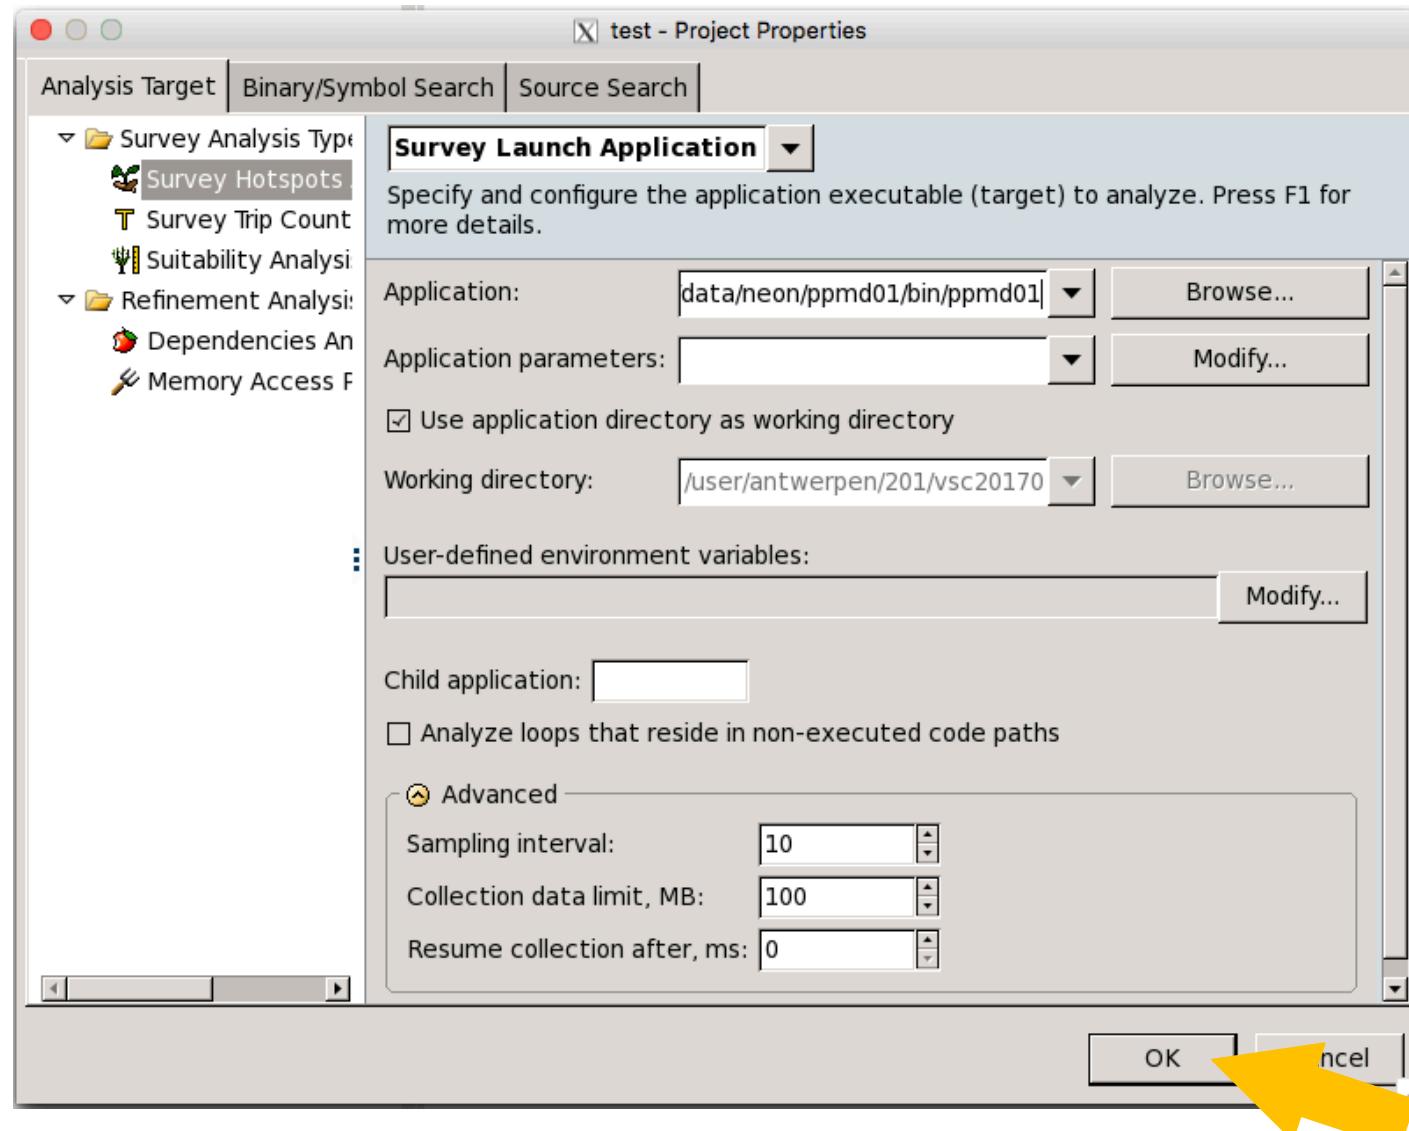
The compute node we are running on

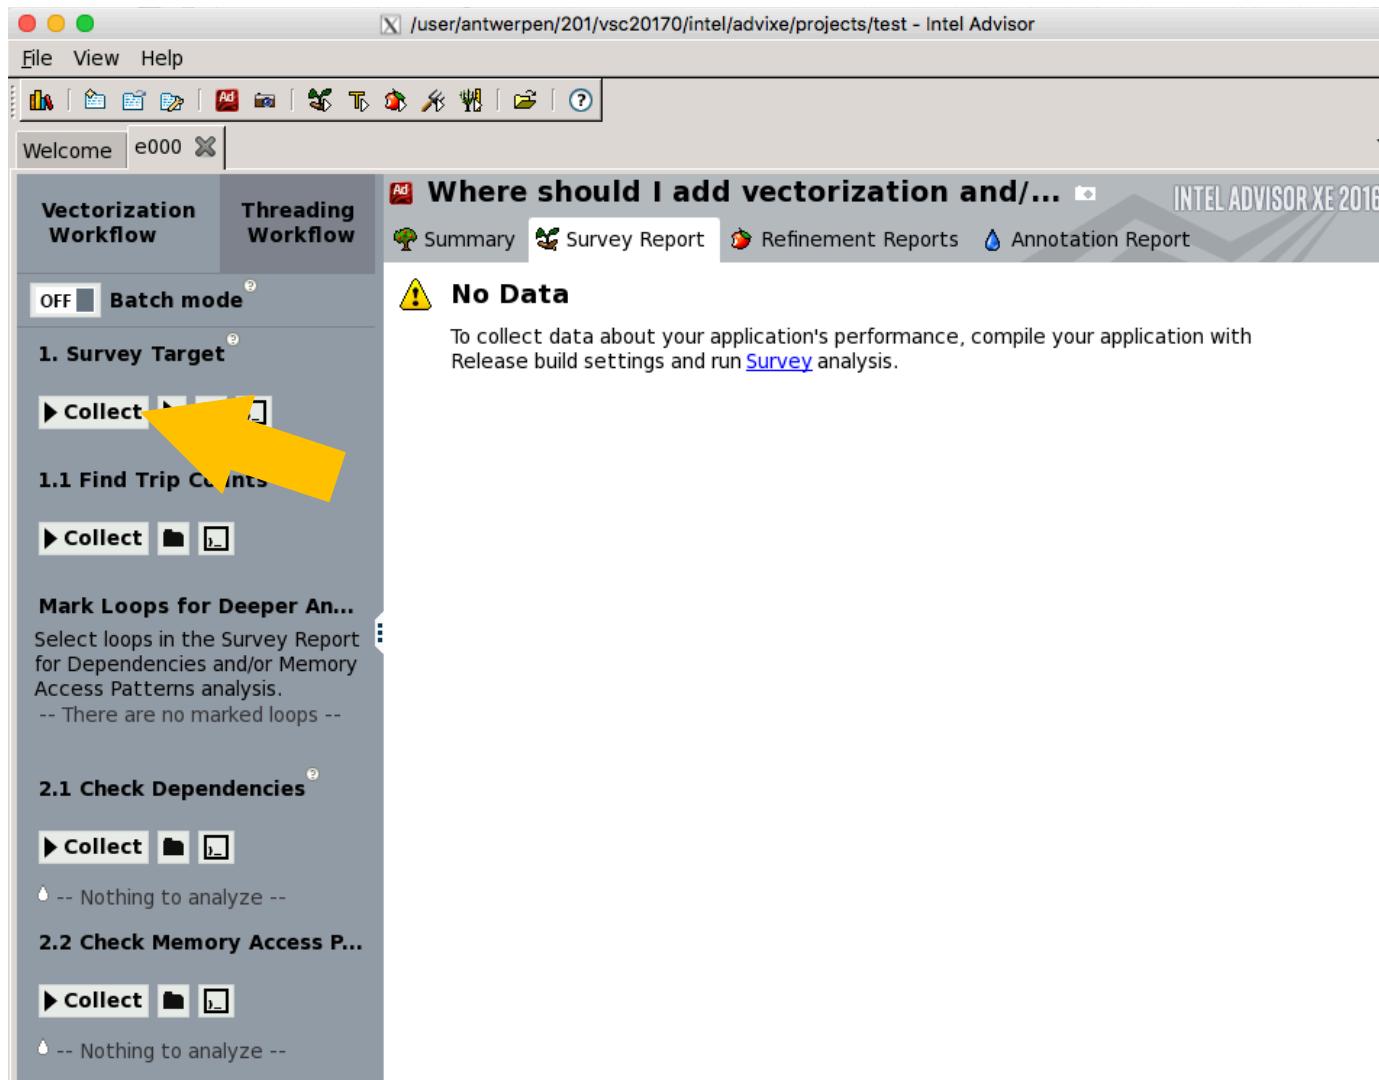
& = Run in background, so the terminal remains functional











Where should I add vectorization and/or threading parallelism? 

Elapsed time: 2162.89s   FILTER: All Module  All Sources  Loop  All Threads

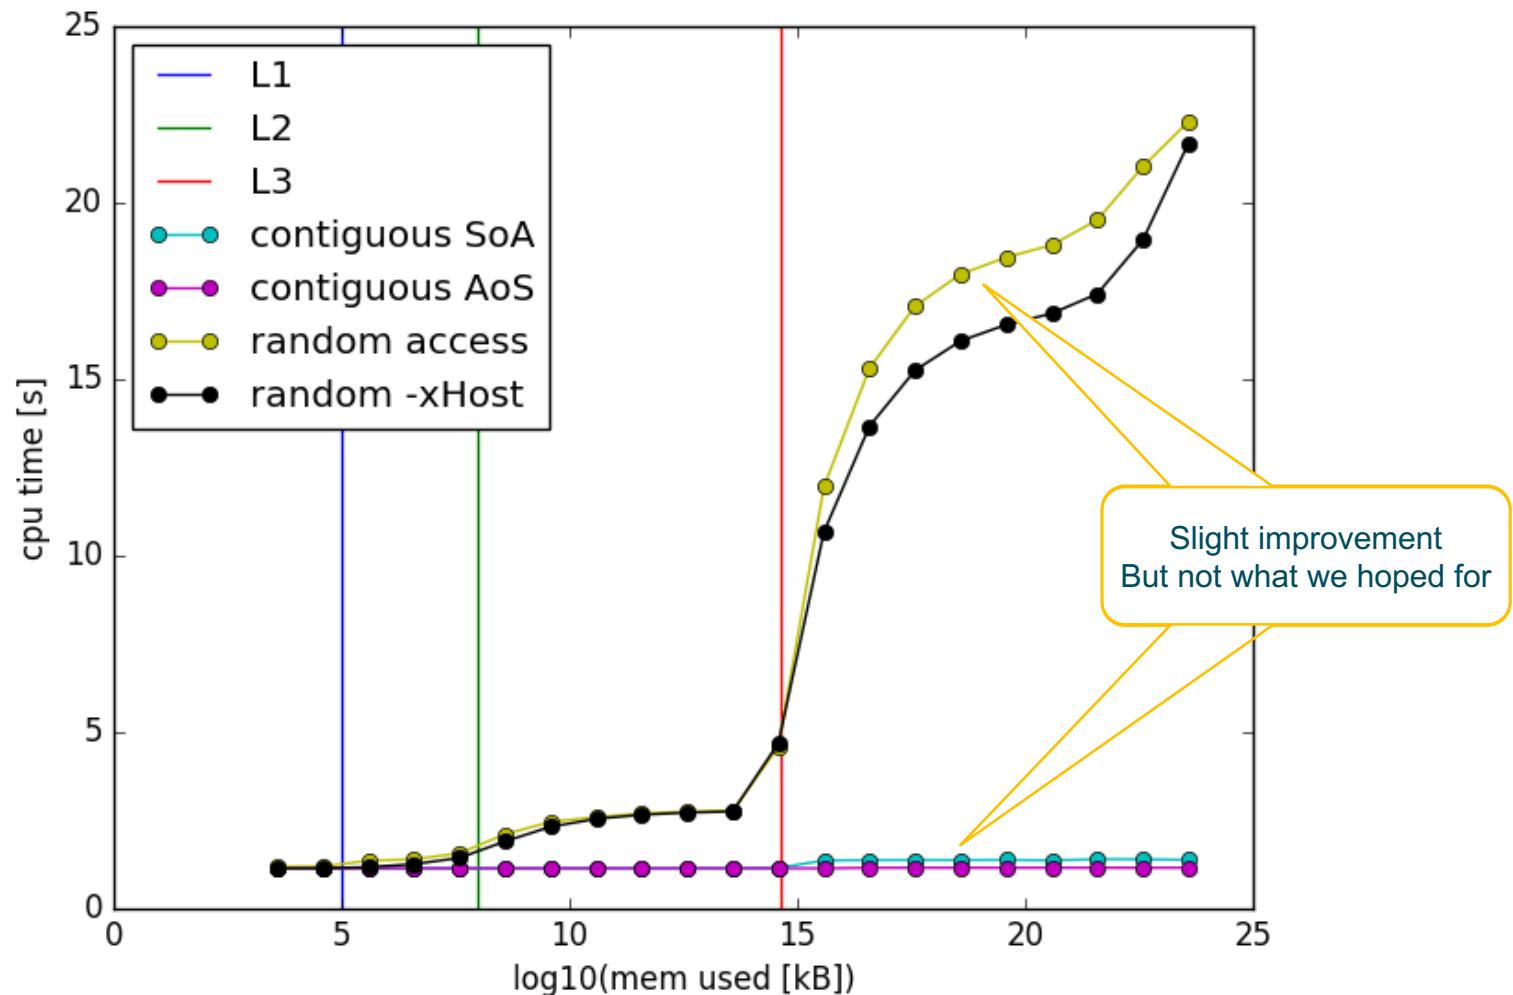
 Summary  Survey Report  Refinement Reports  Annotation Report

 **Higher instruction set architecture (ISA) available**

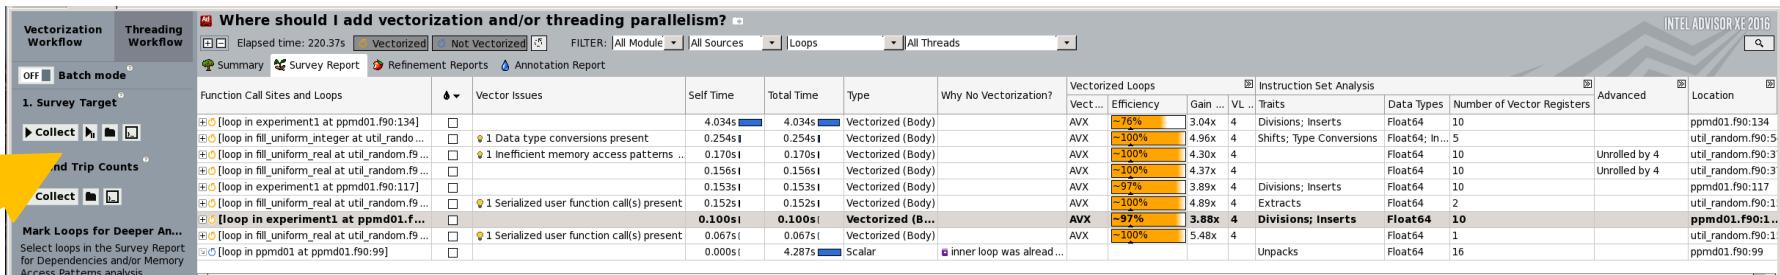
Your application was compiled using an ISA lower than the ISA available on this machine. Consider recompiling your application: on this machine using the highest ISA available - use the `xHost` option, or on the original machine for a higher ISA - use the `x` option.

Function Call Sites and Loops	Vector Issues	Self Time	Total Time	Type	Why No Vectorization?	Vectorized Loops
						Vect... Effic...
 [loop in fill_uniform_integer at util_random.f90:73]	 1 Data type conversion	15.657s	15.657s	Vectorized (Body)		SSE2 ~100%
 [loop in fill_uniform_real at util_random.f90:73]	 1 Inefficient memory access	1.210s	1.210s	Vectorized (Body)		SSE2 ~100%
 [loop in fill_uniform_real at util_random.f90:73]		1.000s	1.000s	Vectorized (Body)		SSE2 ~100%
 [loop in experiment1 at ppmf01.f90:73]	 2 Assumed dependence	0.981s	0.981s	Scalar	 vector dependence	
 [loop in fill_uniform_real at util_random.f90:73]	 1 Serialized usage	0.976s	0.976s	Vectorized (Body)		
 [loop in fill_uniform_real at util_random.f90:73]	 1 Serialized usage	0.359s	0.359s	Vectorized (Body)		SSE ~100%
 [loop in ppmf01 at ppmf01.f90:26]		0.000s	20.183s	Scalar	 loop with function call	

- Add compiler option `-xHost`



- Run Intel Advisor again



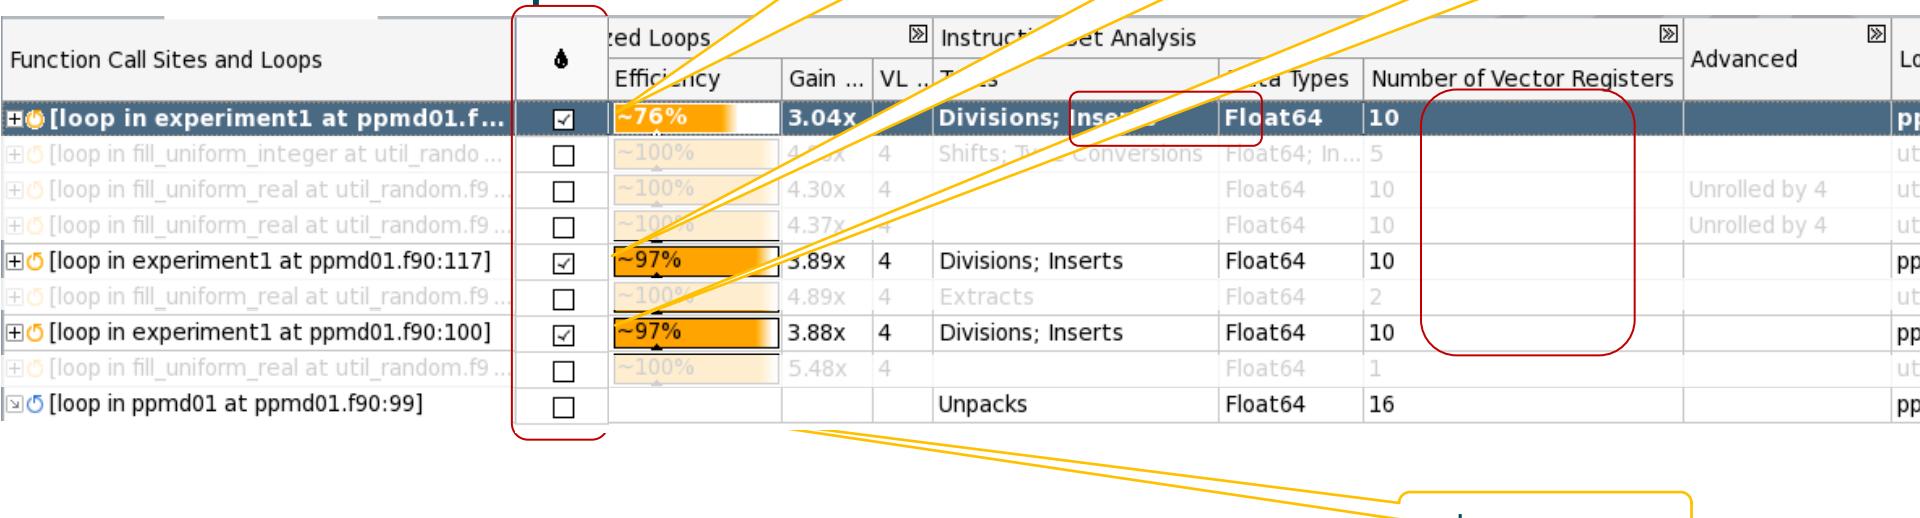
Where should I add vectorization and/or threading parallelism?

Function Call Sites and Loops

	Self Time	Total Time	Type	Why No Vectorization?	Vectorized Loops	Instruction Set Analysis	Advanced	Location
[loop in experiment1 at ppmd01.f90:134]	4.034s	4.034s	Vectorized (Body)		AVX ~76% 3.04x 4 Divisions; Inserts	Float64 10		ppmd01.f90:134
[loop in fill_uniform_integer at util_random.f90:1]	0.254s	0.254s	Vectorized (Body)		AVX ~100% 4.96x 4 Shifts: Type Conversions	Float64; In... 5		util_random.f90:5
[loop in fill_uniform_real at util_random.f90:1]	0.170s	0.170s	Vectorized (Body)		AVX ~100% 4.30x 4	Float64 10	Unrolled by 4	util_random.f90:3
[loop in fill_uniform_real at util_random.f90:1]	0.156s	0.156s	Vectorized (Body)		AVX ~100% 4.37x 4	Float64 10	Unrolled by 4	util_random.f90:3
[loop in experiment1 at ppmd01.f90:117]	0.153s	0.153s	Vectorized (Body)		AVX ~91% 3.89x 4 Divisions; Inserts	Float64 10		ppmd01.f90:117
[loop in fill_uniform_real at util_random.f90:1]	0.152s	0.152s	Vectorized (Body)		AVX ~100% 4.89x 4 Extracts	Float64 2		util_random.f90:1
[loop in experiment1 at ppmd01.f90:1]	0.100s	0.100s	Vectorized (B...)		AVX ~97% 3.88x 4 Divisions; Inserts	Float64 10		ppmd01.f90:1..
[loop in fill_uniform_real at util_random.f90:1]	0.067s	0.067s	Vectorized (Body)		AVX ~100% 5.48x 4	Float64 1		util_random.f90:1
[loop in ppmd01 at ppmd01.f90:99]	0.000s	4.287s	Scalar	inner loop was already...	Unpacks	Float64 16		ppmd01.f90:99

Mark Loops for Deeper An...
Select loops in the Survey Report for Dependencies and/or Memory Access Patterns analysis.

- List of hot spots



Function Call Sites and Loops

	Efficiency	Gain ...	VL ...	Tra... ts	Data Types	Number of Vector Registers	Advanced	Loc...
[loop in experiment1 at ppmd01.f...	~76%	3.04x	Divisions; Inserts		Float64	10		PP...
[loop in fill_uniform_integer at util_random.f90:1]	~100%	4.96x	4	Shifts: Type Conversions	Float64; In... 5			util...
[loop in fill_uniform_real at util_random.f90:1]	~100%	4.30x	4		Float64 10		Unrolled by 4	util...
[loop in fill_uniform_real at util_random.f90:1]	~100%	4.37x	4		Float64 10		Unrolled by 4	util...
[loop in experiment1 at ppmd01.f90:117]	~97%	3.89x	4	Divisions; Inserts	Float64 10			pp...
[loop in fill_uniform_real at util_random.f90:1]	~100%	4.89x	4	Extracts	Float64 2			util...
[loop in experiment1 at ppmd01.f90:100]	~97%	3.88x	4	Divisions; Inserts	Float64	10		pp...
[loop in fill_uniform_real at util_random.f90:1]	~100%	5.48x	4		Float64 1			util...
[loop in ppmd01 at ppmd01.f90:99]				Unpacks	Float64 16			pp...

Random access Contiguous-AoS Contiguous-SoA

Loop over m

File View Help

Welcome e000 X

Vectorization Workflow Threading Workflow

Elapsed time: 224.49s Vectorized Not Vectorized FILTER: All Module ppmd01.f90 Loops All Threads

Summary Survey Report Refinement Reports Annotation Report

Site Location Loop-Carried Dependencies Strides Distribution Access Pattern Site Name

[loop in experiment1 at ppmd01.f90:1 ...]	No information available	50% / 50% / 0%	Mixed strides	loop_site_15
[loop in experiment1 at ppmd01.f90:1 ...]	No information available	50% / 50% / 0%	Mixed strides	loop_site_18
[loop in experiment1 at ppmd01.f90:1 ...]	No information available	50% / 0% / 50%	Mixed strides	loop_site_20

50%:percentage of memory instructions with unit stride or stride 0 accesses
 Unit stride (stride 1) = Instruction accesses memory that consistently changes by one element from iteration to iteration
Uniform stride (stride 0) = Instruction accesses the same memory from iteration to iteration
0%: percentage of memory instructions with fixed or constant non-unit stride accesses
 Constant stride (stride N) = Instruction accesses memory that consistently changes by N Elements from iteration to iteration
 Example: for the double floating point type, stride 4 means the memory address accessed by this instruction increased by 32 bytes, (4*sizeof(double)) with each iteration

50%: percentage of memory instructions with irregular (variable or random) stride accesses
 Irregular stride = Instruction accesses memory addresses that change by an unpredictable number of elements from iteration to iteration
 Typically observed for indirect indexed array accesses, for example, a[index[i]]
- gather (irregular) accesses, detected for v(p)gather* instructions on AVX2 Instruction Set Architecture
- scatter (irregular) accesses, detected for v(p)scatter* instructions on AVX2 Instruction Set Architecture

Memory Access Patterns Report Dependencies Report Recommendations

ID Stride

P3	1
P6	
P12	-43790546; -32607650; -30850839; -29707033; -21902228; -14150377; -10835

- Advisor tells us
 - Vectorization is ok
 - Strided memory access in the random access loop is a problem
- Let's run a **memory access** analysis in VTune

```
$ssh -X vsc20170@login.hpc.uantwerpen.be  
Last login: Thu Sep  8 16:38:25 2016 from 143.169.185.55
```

Welcome to Hopper!

...

```
vsc20170@ln02 ~$ module load VTune
```

```
vsc20170@ln02 ~$ module list
```

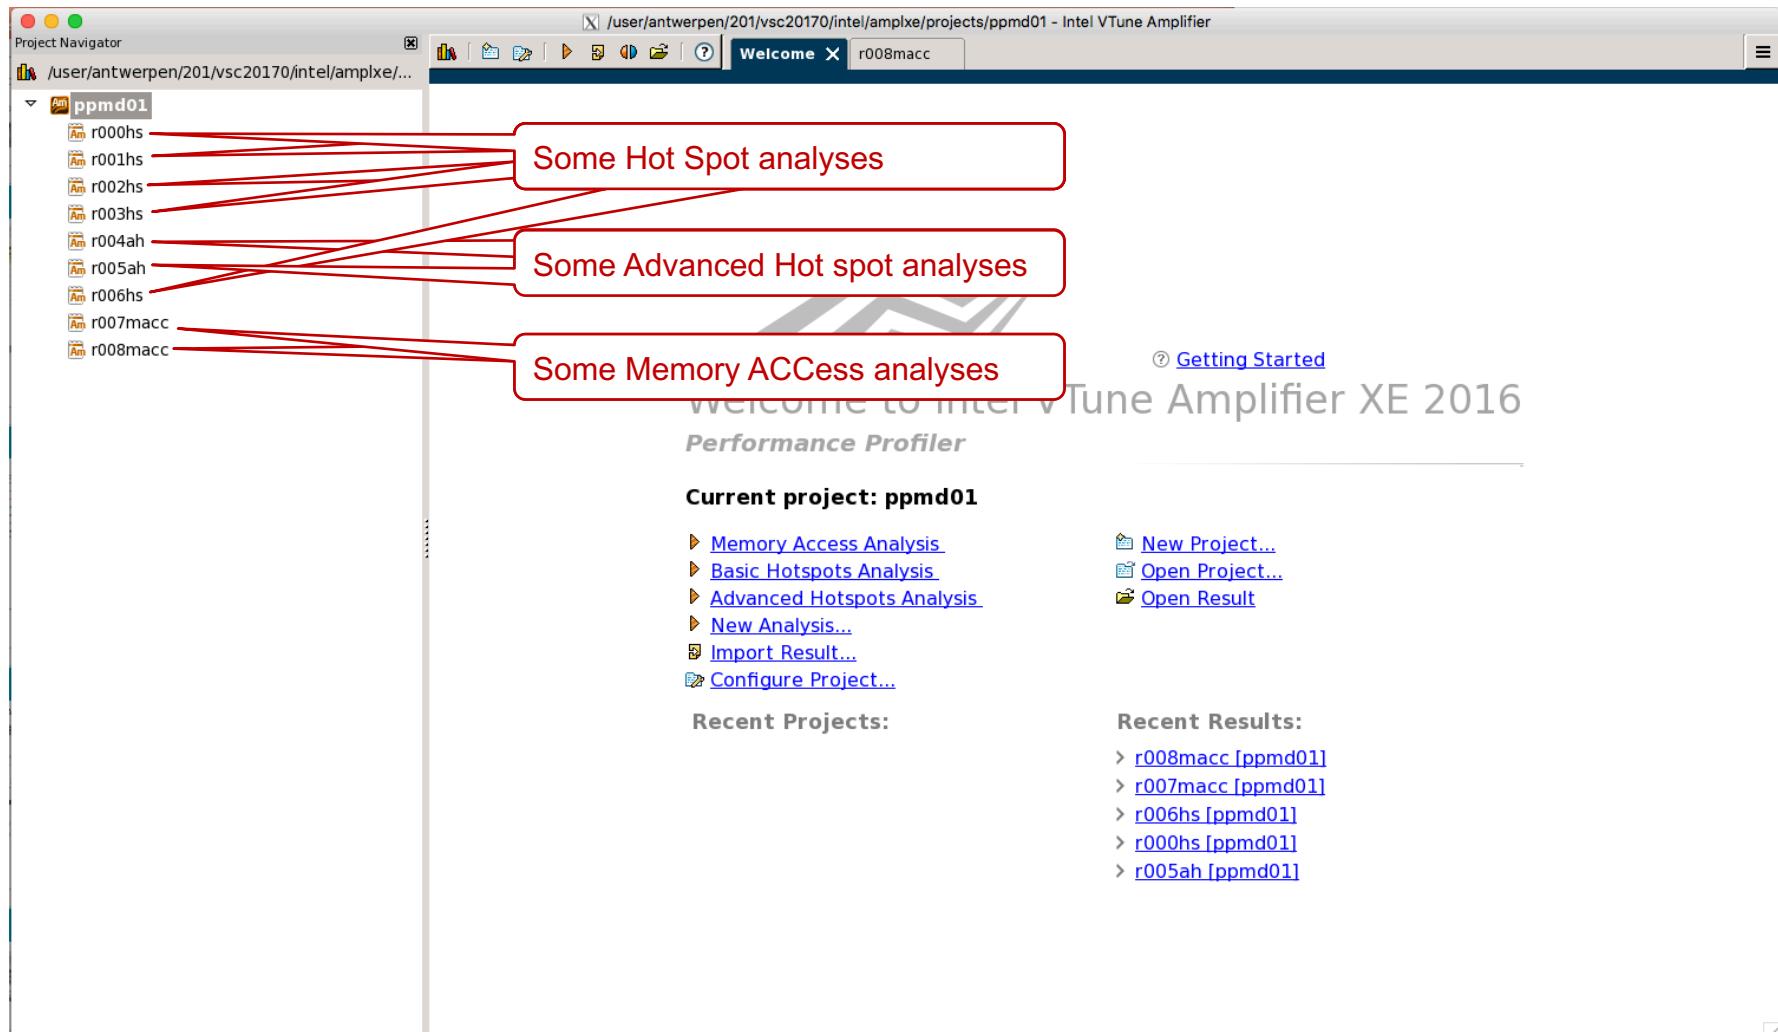
Currently Loaded Modulefiles:

- | | | |
|----------------------------------|---------------------------------------|-----------------------|
| 1) GCCcore/5.4.0 | 4) ifort/2016.3.210-GCC-5.4.0-2.26 | 7) VTune/2016_update3 |
| 2) binutils/2.26-GCCcore-5.4.0 | 5) iccifort/2016.3.210-GCC-5.4.0-2.26 | |
| 3) icc/2016.3.210-GCC-5.4.0-2.26 | 6) Advisor/2016_update4 | |

```
vsc20170@ln02 ~$
```

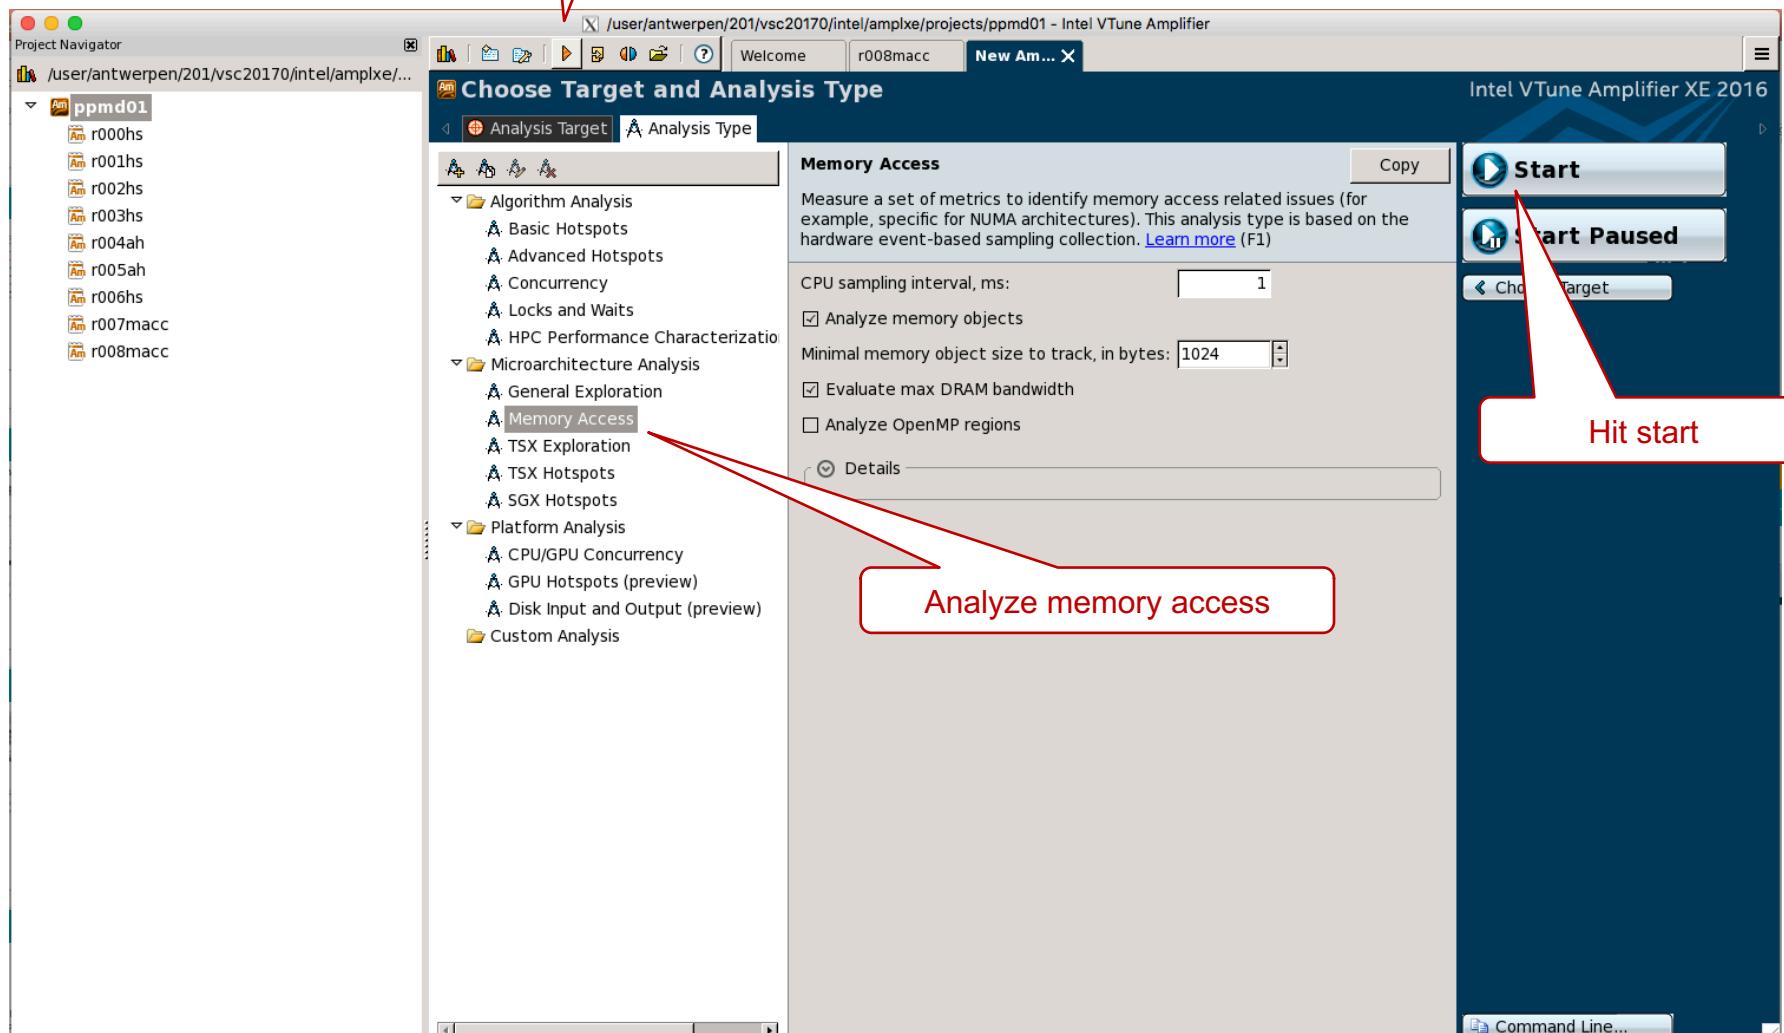
```
vsc20170@ln02 ~$ amplxe-gui &
```

Intel Vtune Amplifier

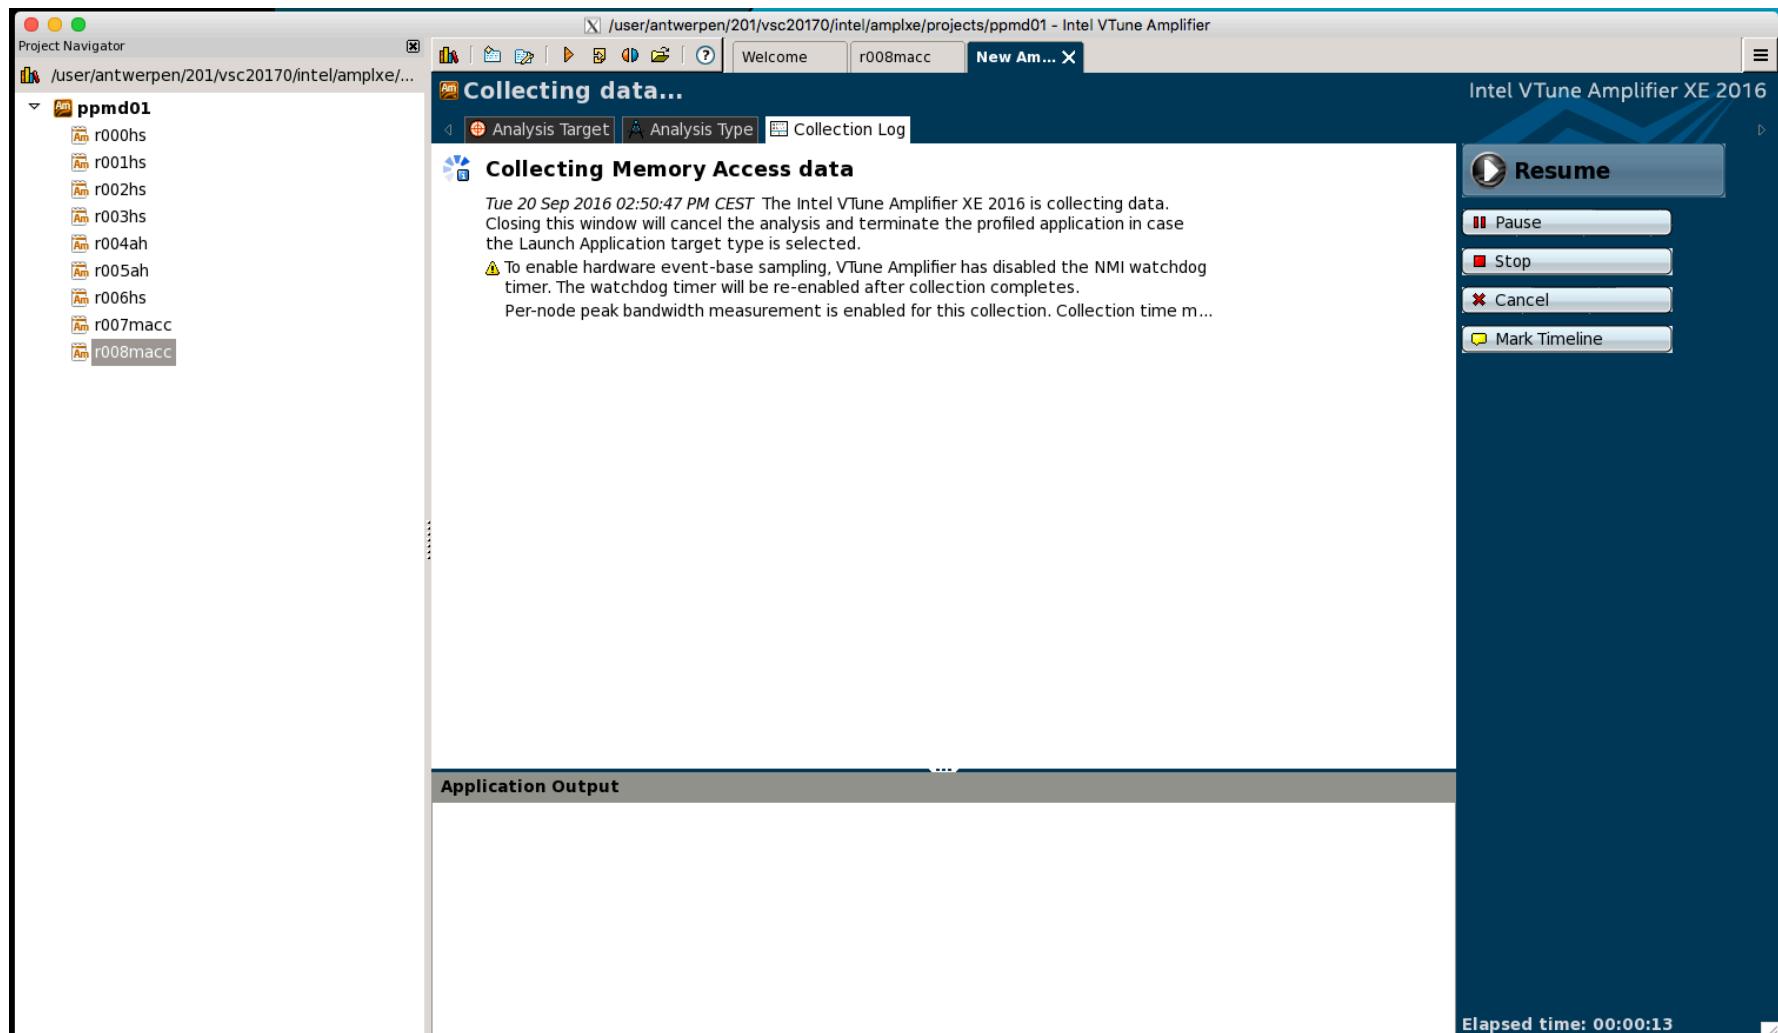


Intel Vtune Amplifier

New analysis



Intel Vtune Amplifier



Intel Vtune Amplifier

Project Navigator

/user/antwerpen/201/vsc20170/intel/amplxe/projects/ppmd01 - Intel VTune Amplifier

Memory Access Memory Usage viewpoint (change) ②

Analysis Target Analysis Type Collection Log Summary Bottom-up Platform

Elapsed Time ②: 44.130s

CPU Time: 42.827s
58.0%

Memory Bound:
The metric value is high. This can indicate that the significant fraction of execution pipeline slots could be stalled due to demand memory load and stores. Use Memory Access analysis to have the metric breakdown by memory hierarchy, memory bandwidth information, correlation by memory objects.

L1 Bound: 0.033
L2 Bound: 0.000
L3 Bound: 0.000

DRAM Bound: 0.470

This metric shows how often CPU was stalled on the main memory (DRAM)

Memory Bandwidth: 0.712

This metric represents a fraction of cycles during which an application does not aggregate requests from other threads/cores/sockets (see Uncore counters for that). Consider improving data locality in NUMA multi-socket systems.

Memory Latency: 0.252

This metric represents a fraction of cycles during which an application could be stalled due to the latency of the main memory (DRAM). This metric does not aggregate requests from other threads/cores/sockets (see Uncore counters for that). Consider optimizing data layout or using Software Prefetches (through the compiler).

Remote / Local DRAM Ratio: 0.000
Local DRAM: 1.000

The number of CPU stalls on loads from the local memory exceeds the threshold. Consider caching data to improve the latency and increase the performance.

Remote DRAM: 0.000
Remote Cache: 0.000

Loads: 44,152,066,228
Stores: 10,512,015,768

LLC Miss Count: 1,600,048,000

Average Latency (cycles): 43
Total Thread Count: 4
Paused Time: 0s

Top Memory Objects

This section lists the most actively used memory objects in your application.

Memory Object	Loads	Stores	LLC Miss Count
[Unknown]	19,448,029,172	[Unknown]	1,600,048,000

Our program is memory bound

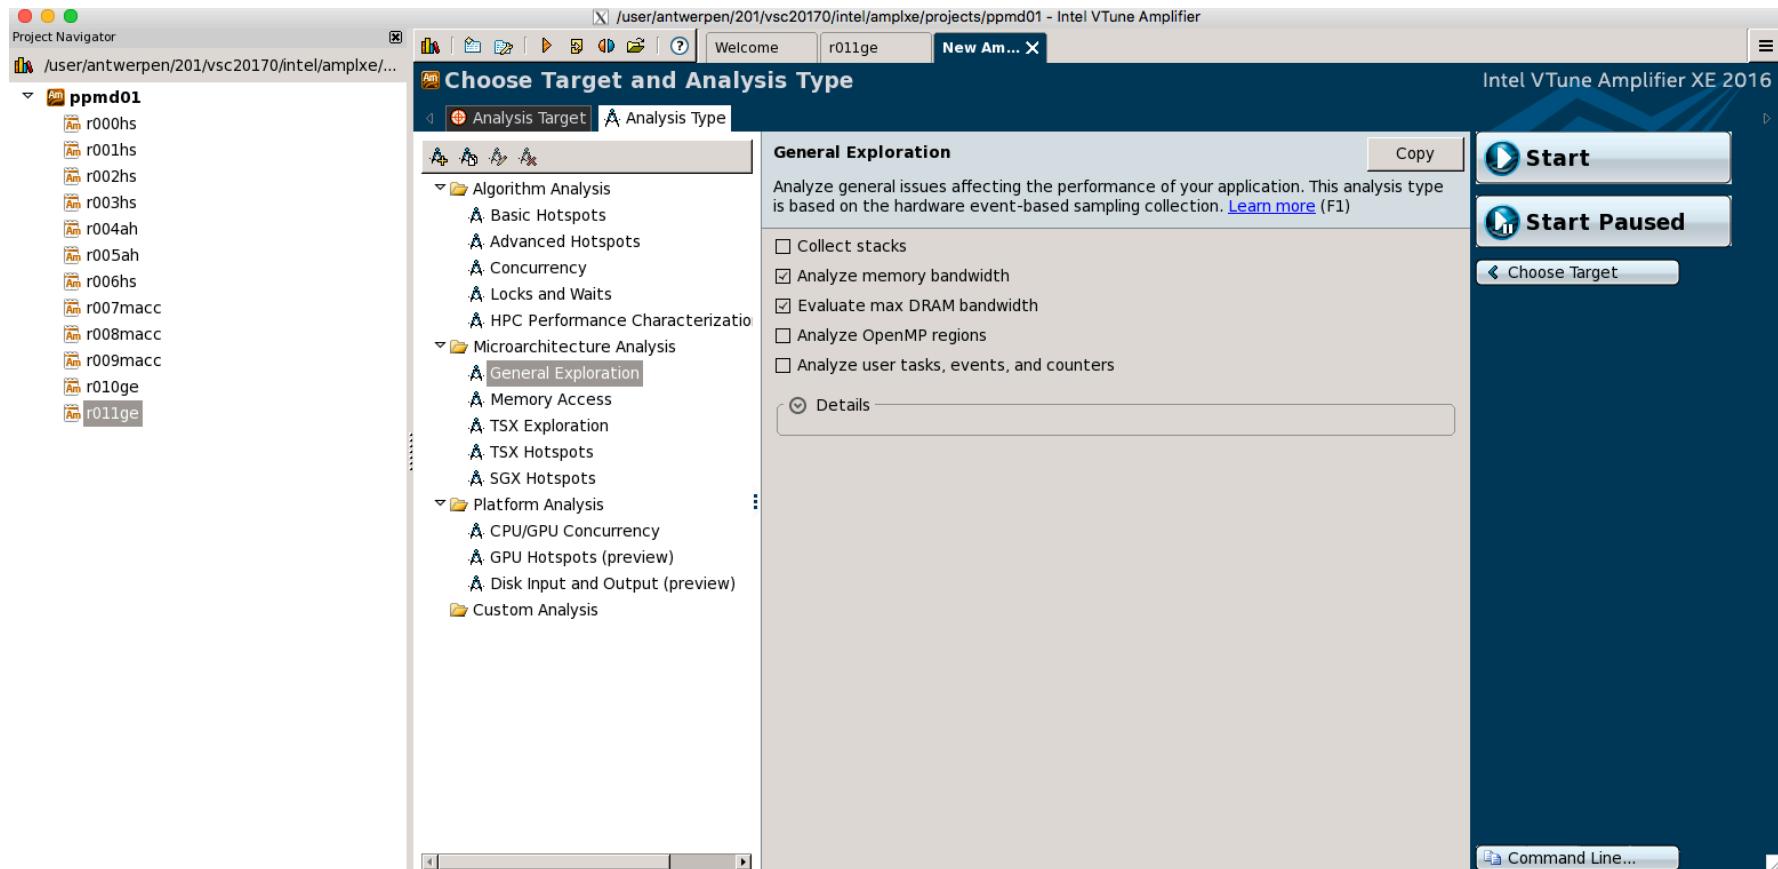
Due to complete cache misses

Bandwidth saturates fast because we move an entire cache line for almost every data item

complete cache misses

Average # of cycles we have to wait for a data item (should be ~1!)

Intel Vtune Amplifier



General Exploration

This view uses [hardware event-based metrics](#) to show code regions that experienced potentially significant architectural bottlenecks. Hover over a metric name in the [grid](#) for the metric description.

Use this view to:

- Identify code regions (modules, functions, and so on) with the highest execution time.
- [Analyze detected hardware issues](#) highlighted by pink cells and get tuning recommendations.

Press **F1** for help on each window.

Show the analysis description when result opens

This metric represents a fraction of cycles during which an application could be stalled due to approaching bandwidth limits of the main memory (DRAM). This metric does not aggregate requests from other threads/cores/sockets (see Uncore counters for that). Consider improving data locality in NUMA multi-socket systems.

[Memory Latency](#) **0.190**

This metric represents a fraction of cycles during which an application could be stalled due to the latency of the main memory (DRAM). This metric does not aggregate requests from other threads/cores/sockets (see Uncore counters for that). Consider optimizing data layout or using Software Prefetches (through the compiler).

Local DRAM	1.000
The number of CPU stalls on loads from the local memory exceeds the threshold. Consider caching data to improve the latency and increase the performance.	
Remote DRAM	0.000

Intel Vtune Amplifier

Elapsed Time [?]: 35.740s

[Clockticks:](#)

[Instructions Retired:](#)

[CPI Rate [?]:](#)

126,124,189,186

48,346,072,519

2.609

The CPI may be too high. This could be caused by issues such as memory stalls, instruction cache misses or branch mispredictions. Explore the other hardware-related metrics to identify what is causing high CPI.

CPI = cycles per instruction

Peak performance corresponds to 4 instructions per cycle in DP vectorized code. Hence CPI should be between 0.25 and 0.5.

[MUX Reliability [?]:](#)

0.996

[Front-End Bound [?]:](#)

0.7%

[Bad Speculation [?]:](#)

0.0%

[Back-End Bound [?]:](#)

87.3%

Identify slots where no uOps are delivered due to a lack of required resources for access. Front-end bound describes a portion of the pipeline where the out-of-order scheduler dispatches ready uOps. Back-end bound uOps get retired according to program order. Stalls due to data-cache misses or stalls due to the overloaded divider unit are examples of back-end bound issues.

[Memory Bound [?]:](#)

66.2%

The metric value is high. This can indicate that the significant fraction of execution time is spent waiting for memory. Use Memory Access analysis to have the metric breakdown by memory hierarchy, memory bandwidth information, correlation by memory objects.

Pipeline stalls (because the data is not arriving in time)

[L1 Bound [?]:](#)

0.040

This metric shows how often machine was stalled without missing the L1 data cache. The L1 cache typically has the shortest latency. However, in certain cases like loads blocked on older stores, a load might suffer a high latency even though it is being satisfied by the L1. Note that this metric value may be highlighted due to DTLB Overhead or Cycles of 1 Port Utilized issues.

[DTLB Overhead [?]:](#)

0.441

A significant proportion of cycles is being spent handling first-level data TLB misses. As with ordinary data caching, focus on improving data locality and reducing working-set size to reduce DTLB overhead. Additionally, consider using profile-guided optimization (PGO) to collocate frequently-used data on the same page. Try using larger page sizes for large amounts of frequently-used data.

Our program is memory bound

[Loads Blocked by Store Forwarding [?]:](#)

0.000

[Lock Latency [?]:](#)

0.000

[Split Loads [?]:](#)

0.000

[4K Aliasing [?]:](#)

0.002

[L2 Bound [?]:](#)

0.000

[L3 Bound [?]:](#)

0.000

[DRAM Bound [?]:](#)

0.576

This metric shows how often CPU was stalled on the main memory (DRAM). Caching typically improves the latency and increases performance.

Our program is memory bound

[Memory Bandwidth [?]:](#)

0.787

This metric represents a fraction of cycles during which an application could be stalled due to memory bandwidth constraints. This metric does not aggregate requests from other threads/cores/sockets (see Uncore counters for that). Consider optimizing data layout or using Software Prefetches (through the compiler).

Bandwidth saturates fast because we move an entire cache line for almost every data item

[Memory Latency [?]:](#)

0.190

This metric represents a fraction of cycles during which an application could be stalled due to the latency of the main memory (DRAM). This metric does not aggregate requests from other threads/cores/sockets (see Uncore counters for that). Consider optimizing data layout or using Software Prefetches (through the compiler).

[Local DRAM [?]:](#)

1.000

The number of CPU stalls on loads from the local memory exceeds the threshold. Consider caching data to improve the latency and increase the performance.

[Remote DRAM [?]:](#)

0.000

[Remote Cache [?]:](#)

0.000

[Store Bound [?]:](#)

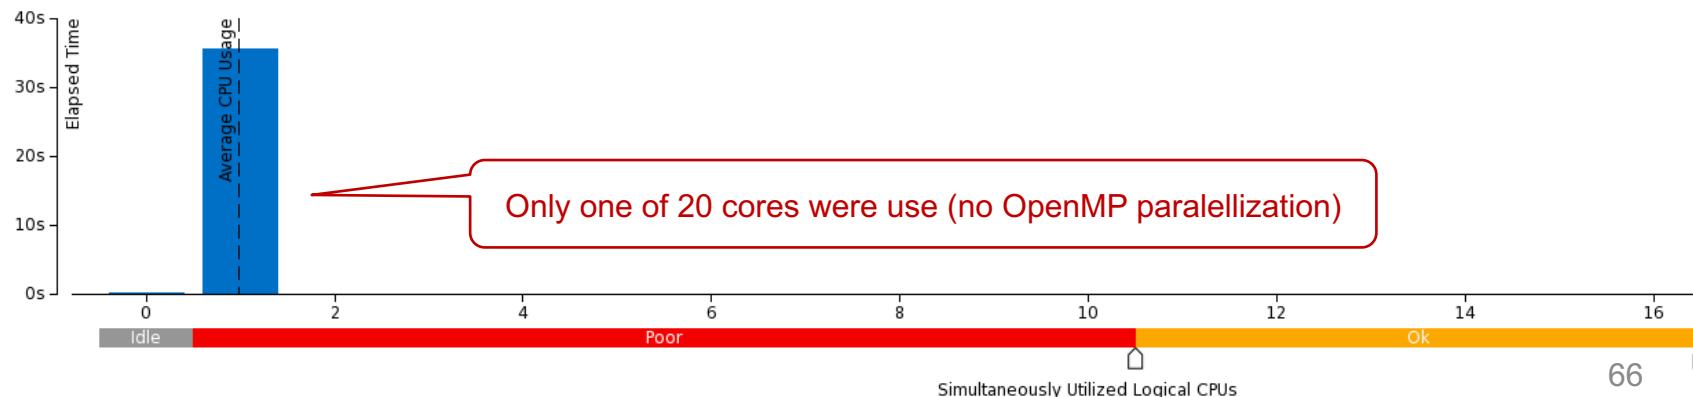
0.134

Intel Vtune Amplifier

Store Bound	0.134
Core Bound	21.1%
This metric represents how much Core non-memory issues were of a bottleneck. Shortage in hardware compute resources, or dependencies software's instructions are both categorized under Core Bound. Hence it may indicate the machine ran out of an OOO resources, certain execution units are overloaded or dependencies in program's data- or instruction- flow are limiting the performance (e.g. FP-chained long-latency arithmetic operations).	
Divider	0.099
Port Utilization	0.373
This metric represents a fraction of cycles during which an application was stalled due to Core non-divider-related issues. For example, heavy data-dependency between nearby instructions, or a sequence of instructions that overloads specific ports. Hint: Loop Vectorization - most compilers feature auto-Vectorization options today - reduces pressure on the execution ports as multiple elements are calculated with same uop.	
Cycles of 0 Ports Utilized	0.627
This metric represents cycles fraction CPU executed no uops on any execution port.	
Cycles of 1 Port Utilized	0.132
Cycles of 2 Ports Utilized	0.089
Cycles of 3+ Ports Utilized	0.068
Retiring	12.0%
Total Thread Count	2
Paused Time	0s

CPU Usage Histogram

This histogram displays a percentage of the wall time the specific number of CPUs were running simultaneously. Spin and Overhead time adds to the Idle CPU usage value.



Advisor vs Vtune amplifier

- Advisor is profiler
 - Analyzes your code on a per statement basis
 - Looks at the assembly code to analyze vectorization
 - Hints to the location of the problem
- Vtune Amplifier accumulates statistics on hardware events such as expensive instructions, vector instructions, cache misses, ...
 - Statistics accumulated on a per subprogram (function, subroutine) basis, not per statement
 - Hints to the nature of the problem
- Both are complementary

What have we learned so far

- Compiler does good job at producing vectorized code
- Advisor will tell you if and why the compiler is sometimes not able to produce vectorized code, and will suggest solutions
- Advisor tells you which parts of your code consume the most cputime and are candidates for optimization

What have we learned so far

- Most often performance problems on modern cpus are due to memory access problems (DRAM latency hits you)
- VTune amplifier gives you clues on how and where to fix the issues
 - CPI and Cache Misses
- Optimize
 1. If there are cache misses, try to reduce them
 - Easier said than done (we'll come to that in the next section)
 2. If you are memory bound and CPI is high,
 1. Verify vectorization
 2. Increase the computational complexity (do more useful work on the data while it is in cache)

Molecular dynamics settings

- Suppose we have 10^9 atoms
- Computing all interactions in single precision
 - $10^9(10^9-1)/2 \sim 0.5 \cdot 10^{18}$
 - complexity $O(N^2)$ – not a good idea
 - Lennard-Jones is short range
 - $\lim_{r \rightarrow \infty} 4\pi r^2 V_{LJ}(r) \rightarrow 0$
 - In practice cut-off $r_c \simeq 2.5$

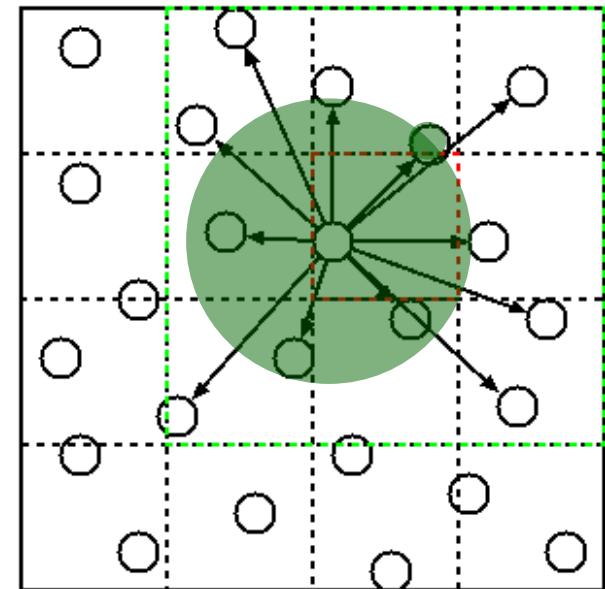
Implementing cut-off

```
forces = 0
do i=1,N
    do j=1,N-1
        r2 = squared_distance(i,j)
        if r2<rcutoff2
            force_ij = ljforce(r2)
            force(i) = force(i) + force_ij
            force(j) = force(j) - force_ij
        endif
    enddo
enddo
integrate forces to update atom positions
```

- Still $O(N^2)$ ☹
- Might be ok for small N

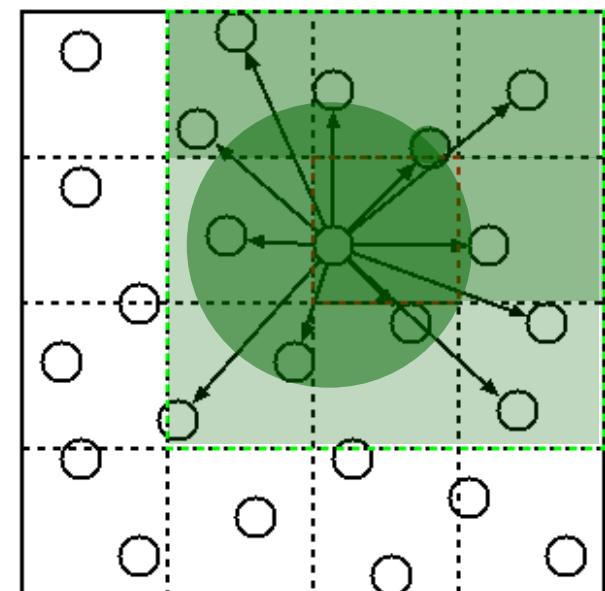
Implementing cut-off

- Verlet lists
 - Verlet list of atom i is list of all atoms j for which $j < i$ and $r_{ij} < r_c$
 - Increase cutoff slightly so that we do not have to update the Verlet lists at every timestep (depending on how vigorously the atoms move)
 - Verlet list construction is amortized
- Construction of Verlet lists is still $O(N^2) \otimes$
- Is dominant data structure: typically between 50 and 100 neighbour atoms/atom



Implementing cut-off

- Put atoms in cells of width r_c : $O(N) \odot$
- Only atoms in neighbouring cells can satisfy $r_{ij} < r_c$
- Because of symmetry only half of the neighbouring cells must be examined
- Construct Verlet lists as follows
 - Loop over all cells $[O(N)]$
 - Loop over all neighbours of the current cell using the neighbour stencil  $[O(1)]$
 - Construct the Verlet list of all atoms in the current cell $[O(1)]$
- Now our MD algorithm is $O(N) \odot$

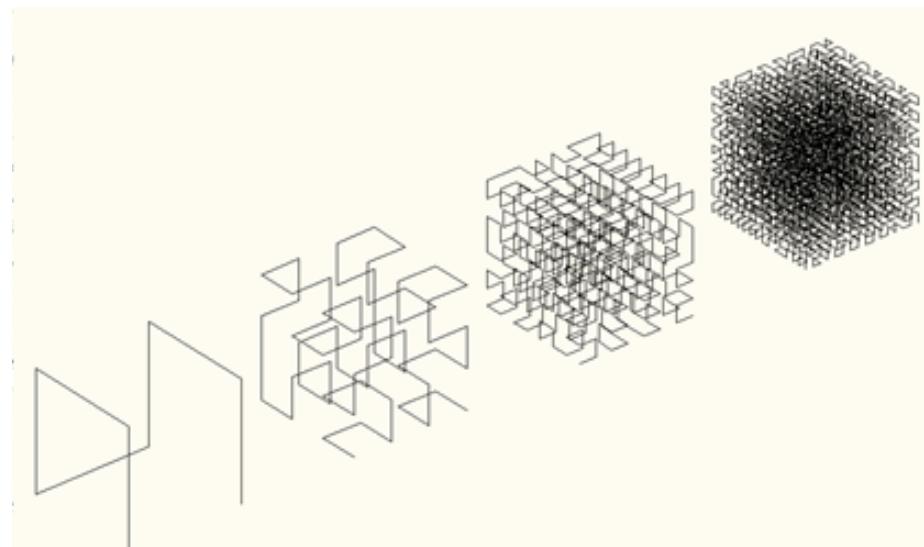


Implementing cut-off

- Atoms move!
- Iterating over the Verlet lists to compute the interactions will soon jump randomly through memory
- Performance evolves naturally to the random access case
- Fix data access pattern using spatial sorting
 - Spatial sort = ensure that atoms which are close in space are also close in memory
 - This reduces cache misses

Fixing the data access pattern

- **Space filling curve**
- Linearize a space of dimension >1
- Hilbert curve
- Hilbert index:
coordinate of a cell along the Hilbert curve
- Locality guarantee:
points close in space are also close along the space filling curve (on average)



Fixing the data access pattern

- 1. Sort atom property arrays ($rx, ry, rz, vx, vy, vz, \dots$) based on the Hilbert index h of the cell of the atoms (**spatial sort**).
Atoms which are close in space (and hence will interact) will be close in memory (and hence will be in the cache with high probability)

- 2. Build a table containing the index of the first atom in each cell, and the number of atoms in the cell (Hilbert list)
- 3. Build Verlet list from the Hilbert list (discard the latter)
- 4. Compute the interactions by looping over the Verlet list
- 5. Integrate forces, updating velocities and positions and time

- 6. If *need_to_rebuild_verlet_list* is true
 jump back to step 1.
else
 continue at step 4.

Fixing the data access pattern

- We need to
 - Compute Hilbert indices
 - Sort atom property arrays
 - Build Hilbert list and Verlet list
- Fixing data access patterns can be a lot of work

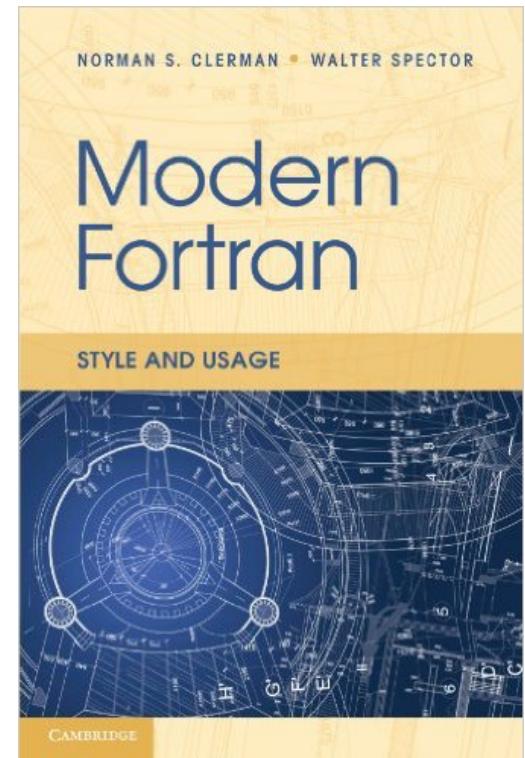
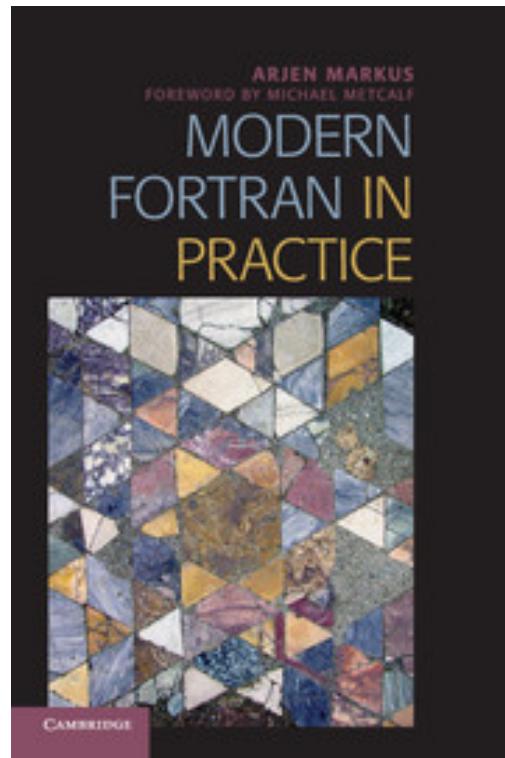
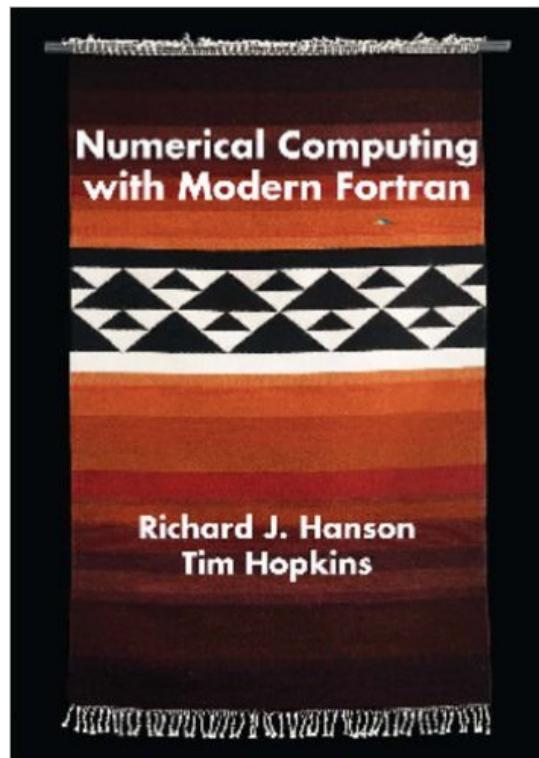
Choosing a programming language

- **Implementation in Fortran? C? C++?**
- Arguments
 - ~~C++ is inefficient~~ Lie #1
 - Modern compilers good enough to generate efficient code
 - After all you are using the same hardware
 - ~~Fortran is efficient~~ Lie #2
 - Also fortran has constructs that sometimes come in handy, but can kill performance
 - But C++ has quite a bit more features which can kill performance than Fortran
 - Because C++ is a general purpose language and Fortran is meant for scientific computing
 - Yet these features can be extremely useful if you use them wisely
 - For computational kernels where performance is an issue you generally need to stay close to the C subset and far away from the C++ features such as classes, inheritance, virtual functions, etc. (templates are an exception)

Choosing a programming language

- ~~I'll use C++ because I know it better~~ Most often a lie too!
 - Unless you have read and understood all the C++ books by Scott Meyers, Herb Sutter, Andrei Alexandrescu, Nicolai Josuttis
 - In which case you probably also understand which C++ features can kill performance and when they should be used to your advantage
 - For number-crunching I find myself advancing faster using Fortran than using C++ (which I do know better!)
- I'll use C++ because it is better documented Not a lie
 - There aren't too many books on Fortran like the above ones on C++
 - There is no website of the same quality as cplusplus.com or cppreference.com for Fortran (imho)
 - But still it is much harder to learn and to learn to use efficiently
 - Not a valid argument

Some useful Fortran books



Choosing a programming language

- I'll use C because that is the language in which Python was written and I want too integrate my code with Python Good point!
- Python integration leverages your code with
 - A high level programming interface:
 - Providing initial data for your simulation is much easier and flexible through a Python script than having to parse input files...
 - Compose and customize high level solution schemes with ease (e.g. choosing another solver for a subproblem)
 - Hundreds of very useful open source Python libraries: Numpy, Scipy, Pandas, matplotlib, ...
- But ...

Choosing a programming language

- I'll use C because that is the language in which Python was written and I want too integrate my code with Python [continued]
 - ...
 - But the easiest way to create your own module that can be imported in Python is through **f2py**
 - Automatically turns your Fortran code into a Python module
 - As simple as
 - `F2py -c mysource.f90 -m myPythonModule`
 - Automatically integrates with **Numpy**! Pass Numpy arrays to your own Fortran library with no effort and no copying of data!
 - Much harder in C
 - Also feasible in C++ with the help of `boost.python` and `boost.multi_array`, easier than C but not as easy as f2py
- ∴ Stick to Fortran

Choosing a programming language

- I'll use C/C++ because I don't want to mess with storage orders
 - Fortran uses row-major ordering indices start at 1
 - C/C++ use column-major ordering indices start at 0
 - Plenty of ways to mess up!
 - Inadvertent copying of the array when passing to fortran!
 - Trivial for 1D arrays (but mind the indexing)
 - Numpy arrays by default use the C convention, but arrays can be easily made to follow the Fortran convention:
 - `A=numpy.empty((2,2), dtype=np.float32, order='F')`
 - A little bit of experimentation will take away the confusion
 - Simplest way to avoid problems:
 - If you use fortran modules adhere to fortran convention in numpy arrays
 - If you use C/C++ modules adhere to C/C++ convention in numpy arrays
 - If you use both pay attention..

Good point!

Choosing a programming language

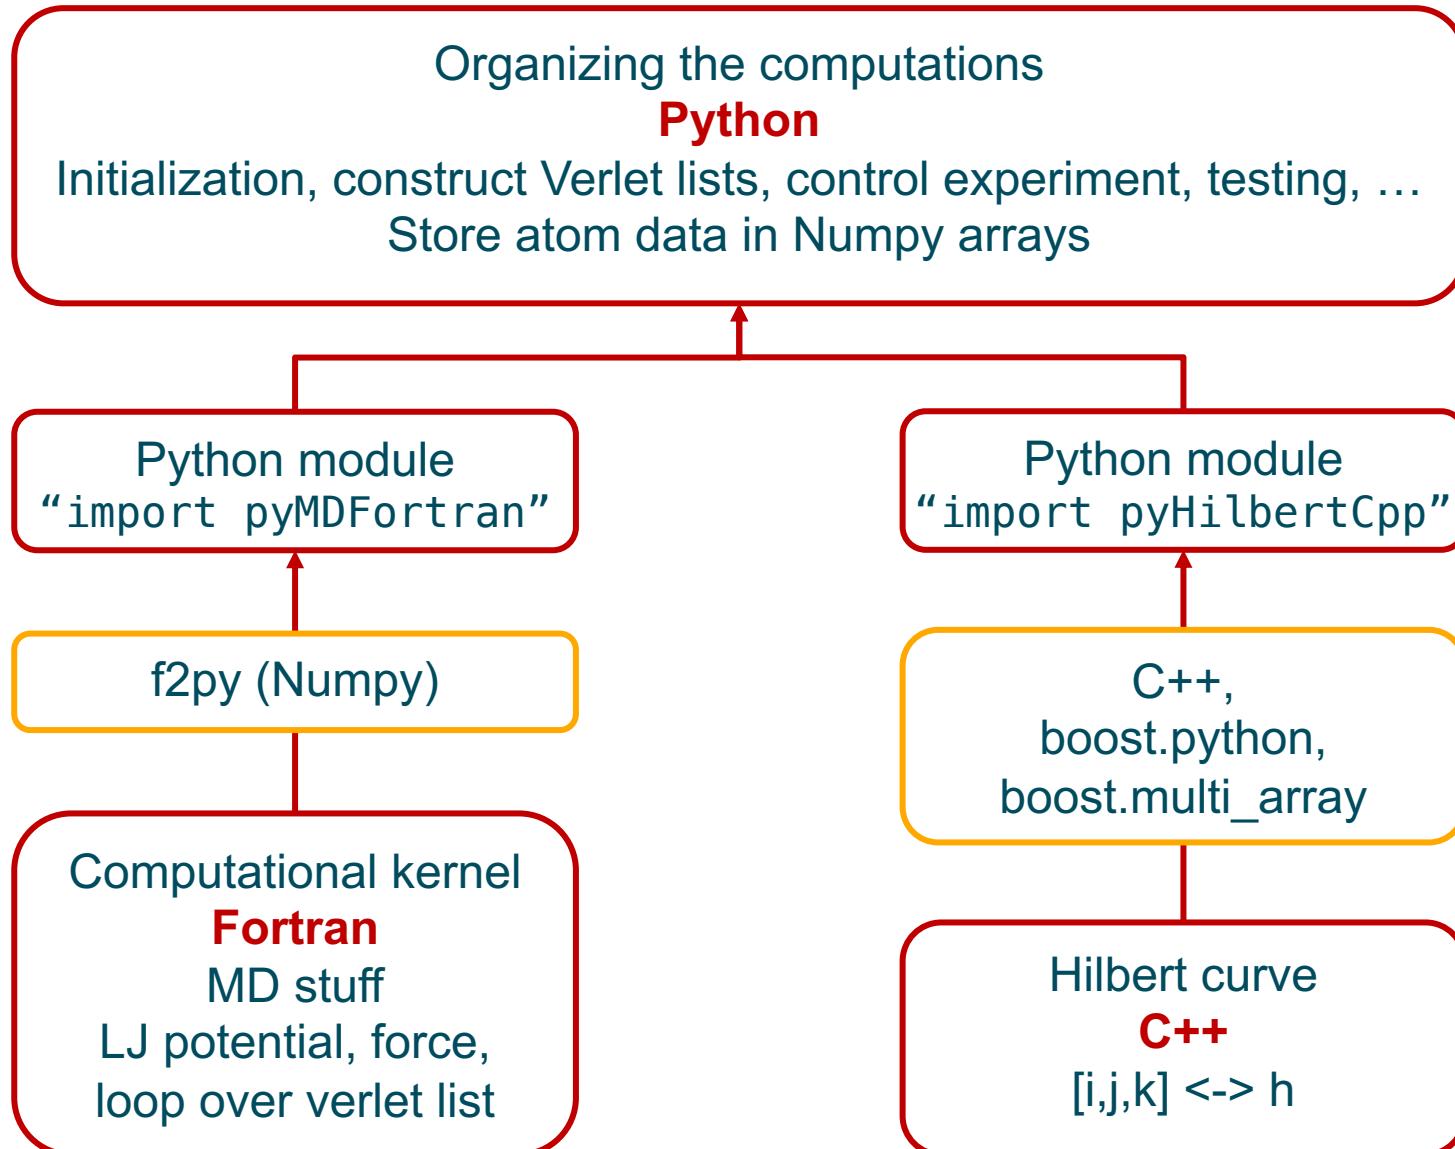
- I'll use X because I need to use some library Y that is written in X
 - Point taken, often the easiest way
 - However, there is a lot of support for mixed language programming
 - If you do not master X you might end up writing inefficient code and loosing the advantage of using Y
 - it may be worthwhile to find out how to tackle the mixed language challenge
 - Once done, you will proceed faster and write efficient code

- Stick to Fortran
- Using Python and Numpy for high level programming and Fortran for your own number-crunching routines is a very practical approach
 - Use f2py to turn your fortran routines into a Python module that is compatible with Numpy

Choosing a programming language

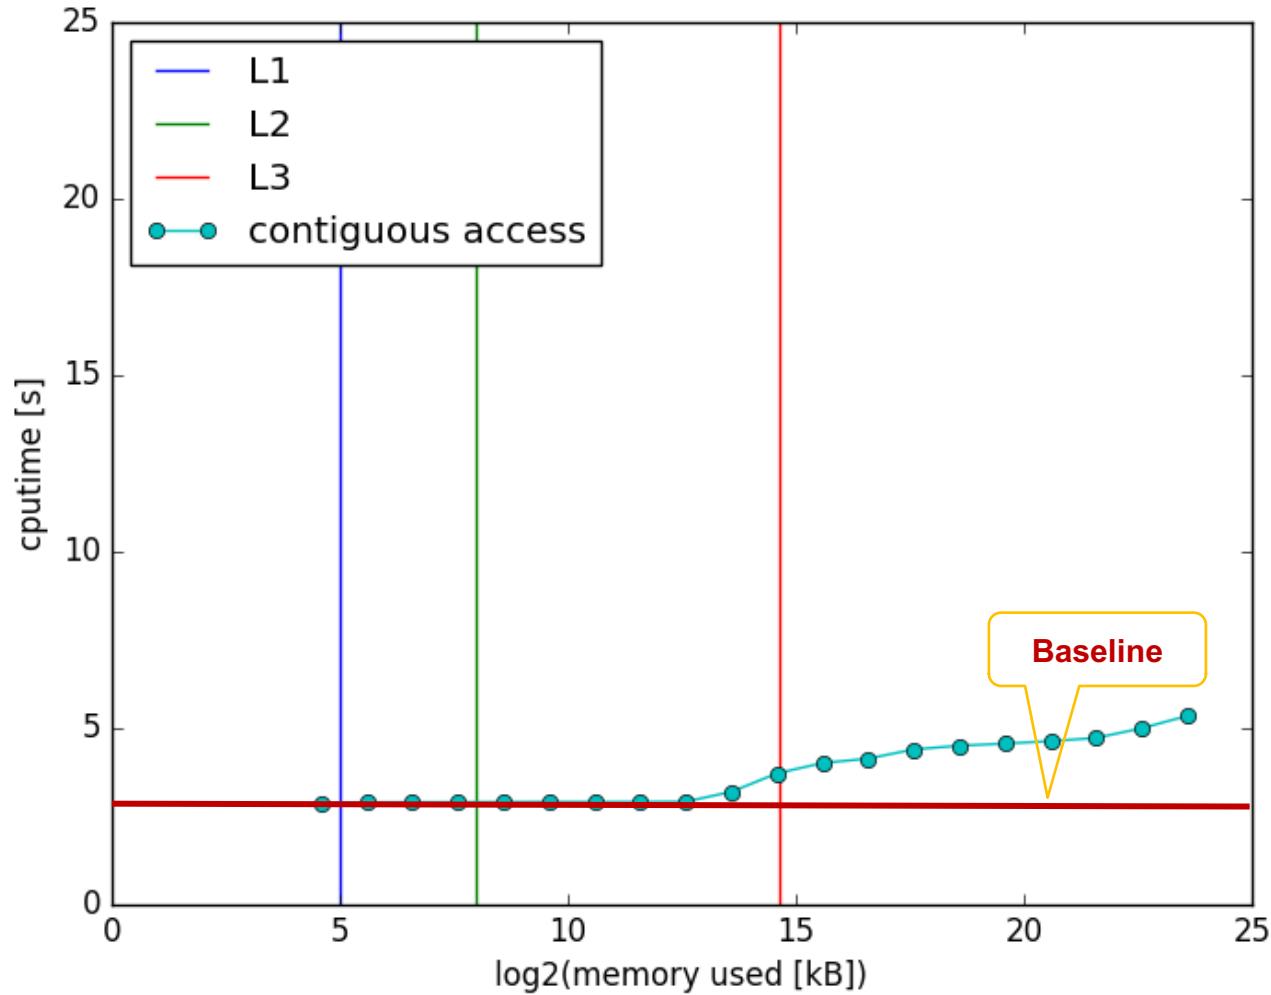
- Details of Python+Numpy+Fortran/C++ is topic of another talk
- Including shared memory parallelism (multi-threading) and distributed memory parallelism (multi-node)

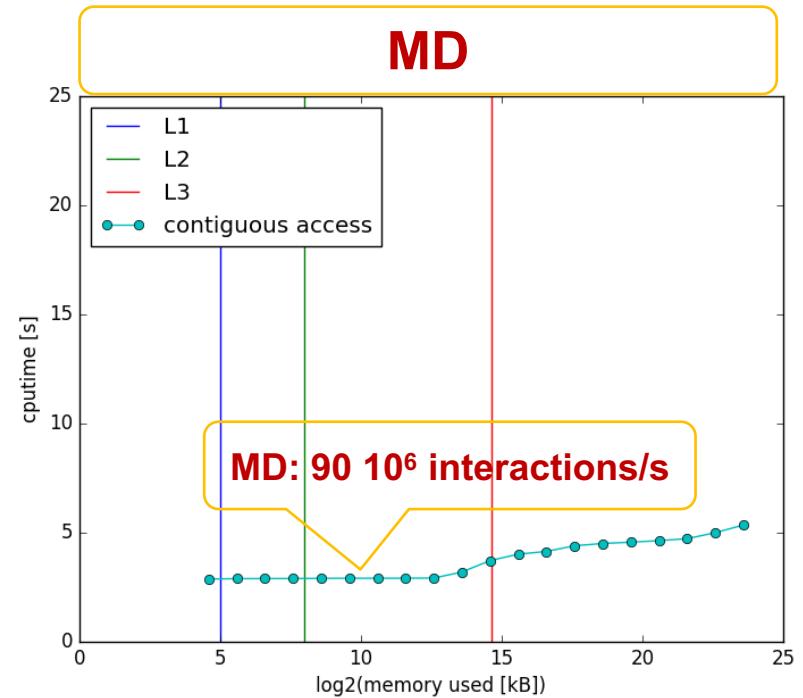
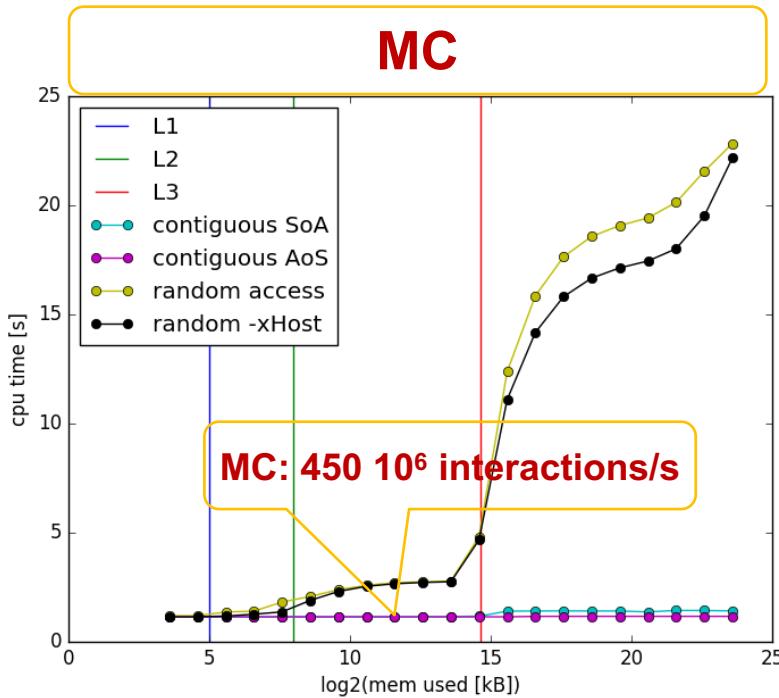
Back to Lennard-Jones MD



- Take derivative of (Lennard-Jones) potential with respect to interatomic distance vector = force exerted on the atoms
- $\frac{dV(r)}{d\vec{r}} = \frac{dV_2(r^2)}{dr^2} \frac{dr^2}{d\vec{r}} = \frac{dV_2(r^2)}{dr^2} 2\vec{r} = f(r^2)\vec{r}$
- Loop over i
 - Loop over $j \in VL_i$
 - $\vec{a}_i += f(r_{ij}^2)\vec{r}_{ij}$
 - $\vec{a}_j -= f(r_{ij}^2)\vec{r}_{ij}$
 - 3 x load (\vec{r}_j)
 - 3 x load (\vec{a}_j)
 - 3 x store (\vec{a}_j)

- Baseline case:
 - N atoms
 - Compute interaction forces of atom 0 with all other atoms
 - Contiguous memory access
 - Bandwidth saturated



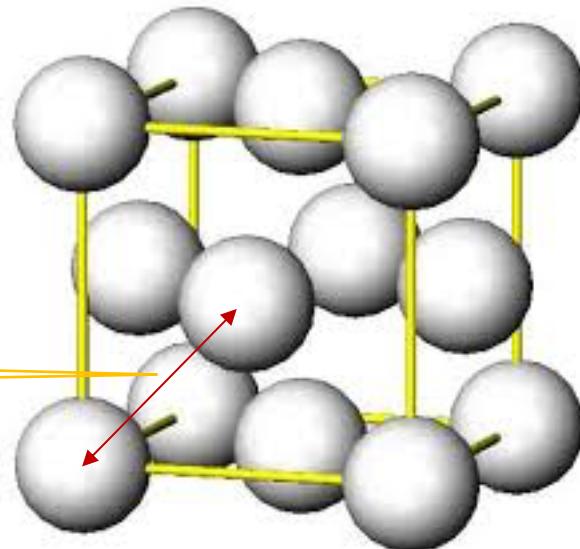


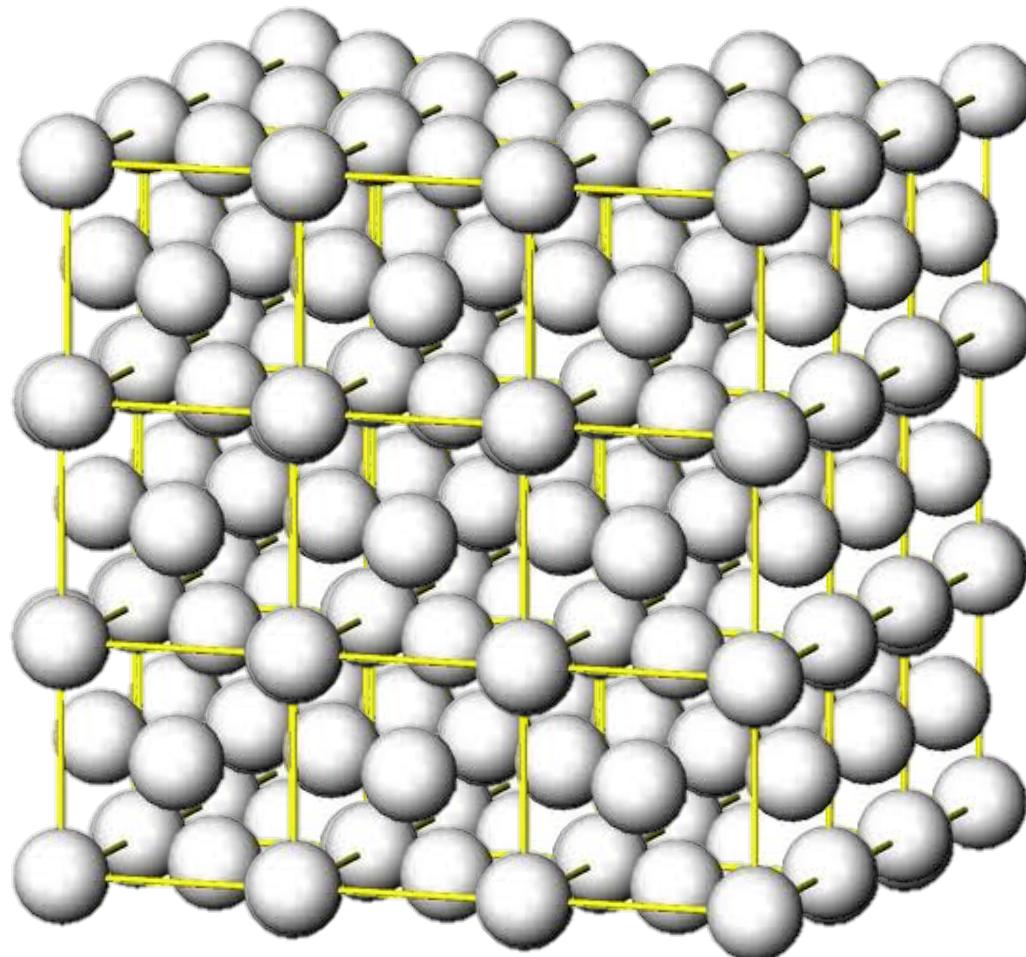
In terms of interactions/s MD is about 5 times slower than MC

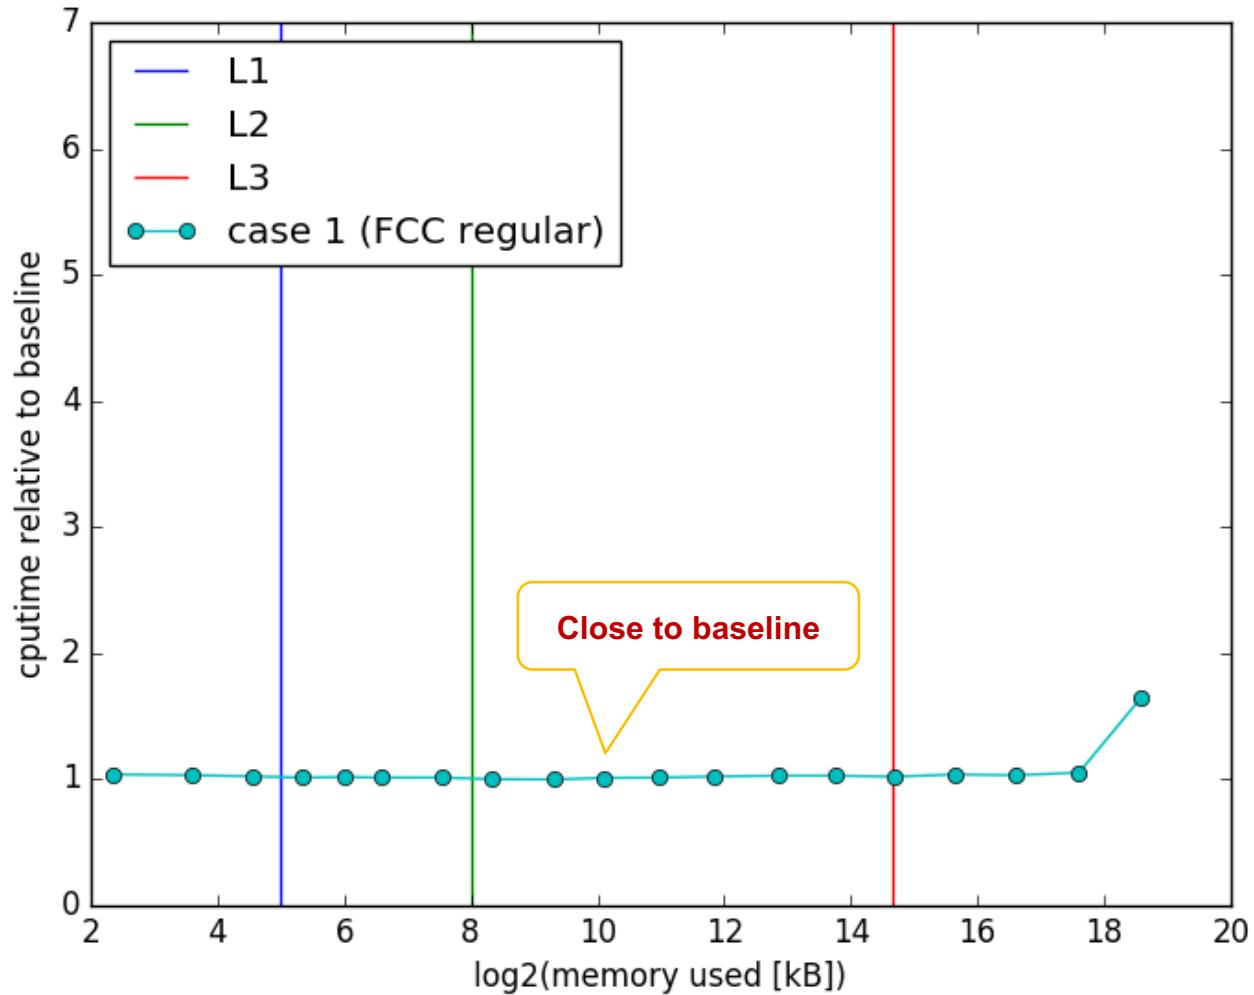
- A bit more instructions per interaction, but MC is memory bound, that should not matter
- 3 times more memory access
- With a read:write ratio of 2:1 the bandwidth drops from 11 GB/s to 9.5 GB/s
- $3 \times 11 / 9.5 = 3.47$
- Still factor 1.44 slower than expected

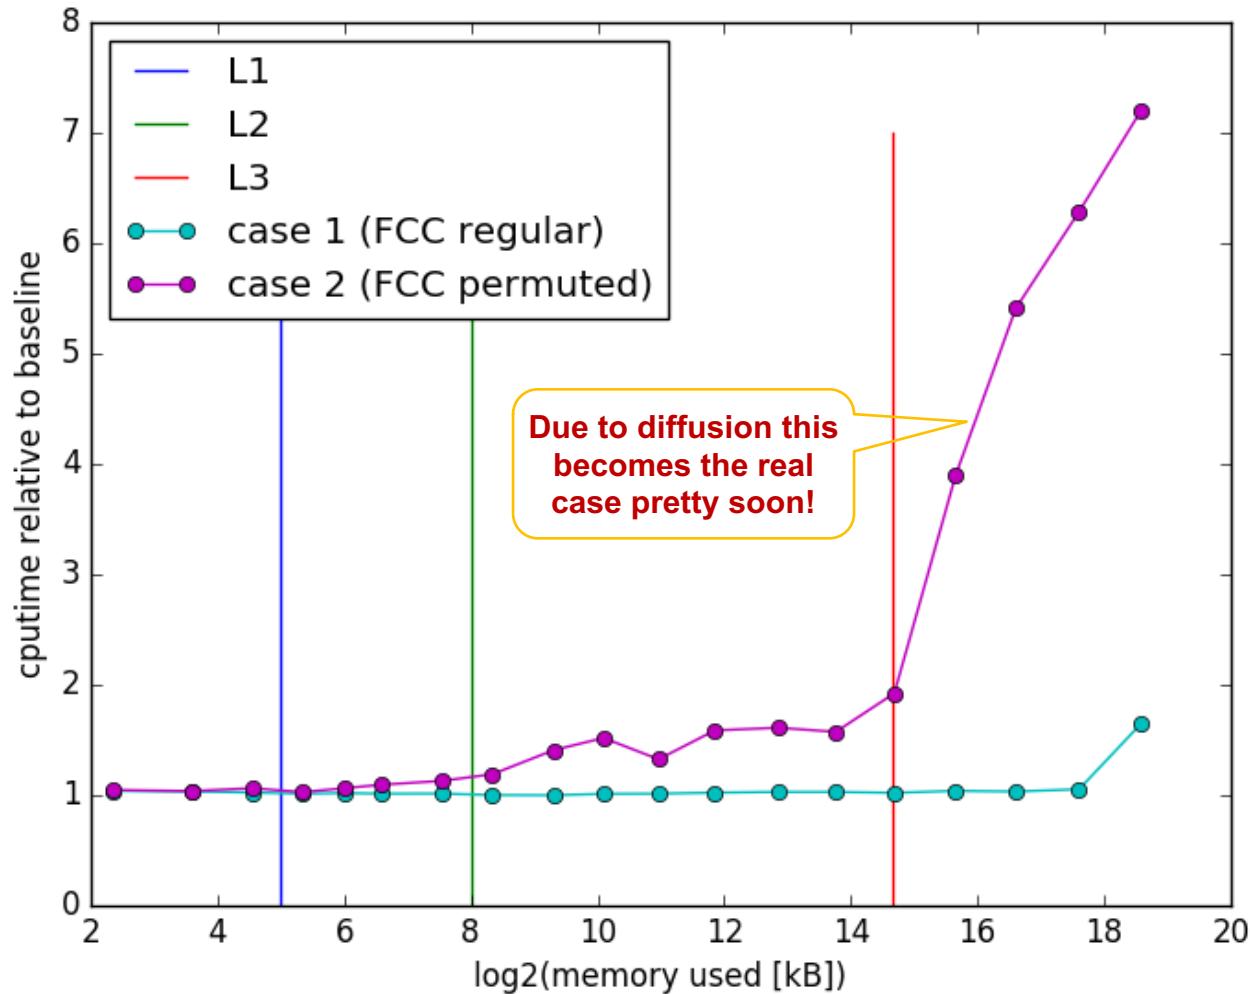
- Three cases:
 1. Atoms on FCC lattice
 2. Permute the atoms (=random memory access)
 3. Spatial sort by hilbert index
- Every experiment build the Verlet list and computes the interactions
- Cputime is measured only for computing the interactions
- Plot result relative to baseline
 - $90 \cdot 10^6$ interactions/s

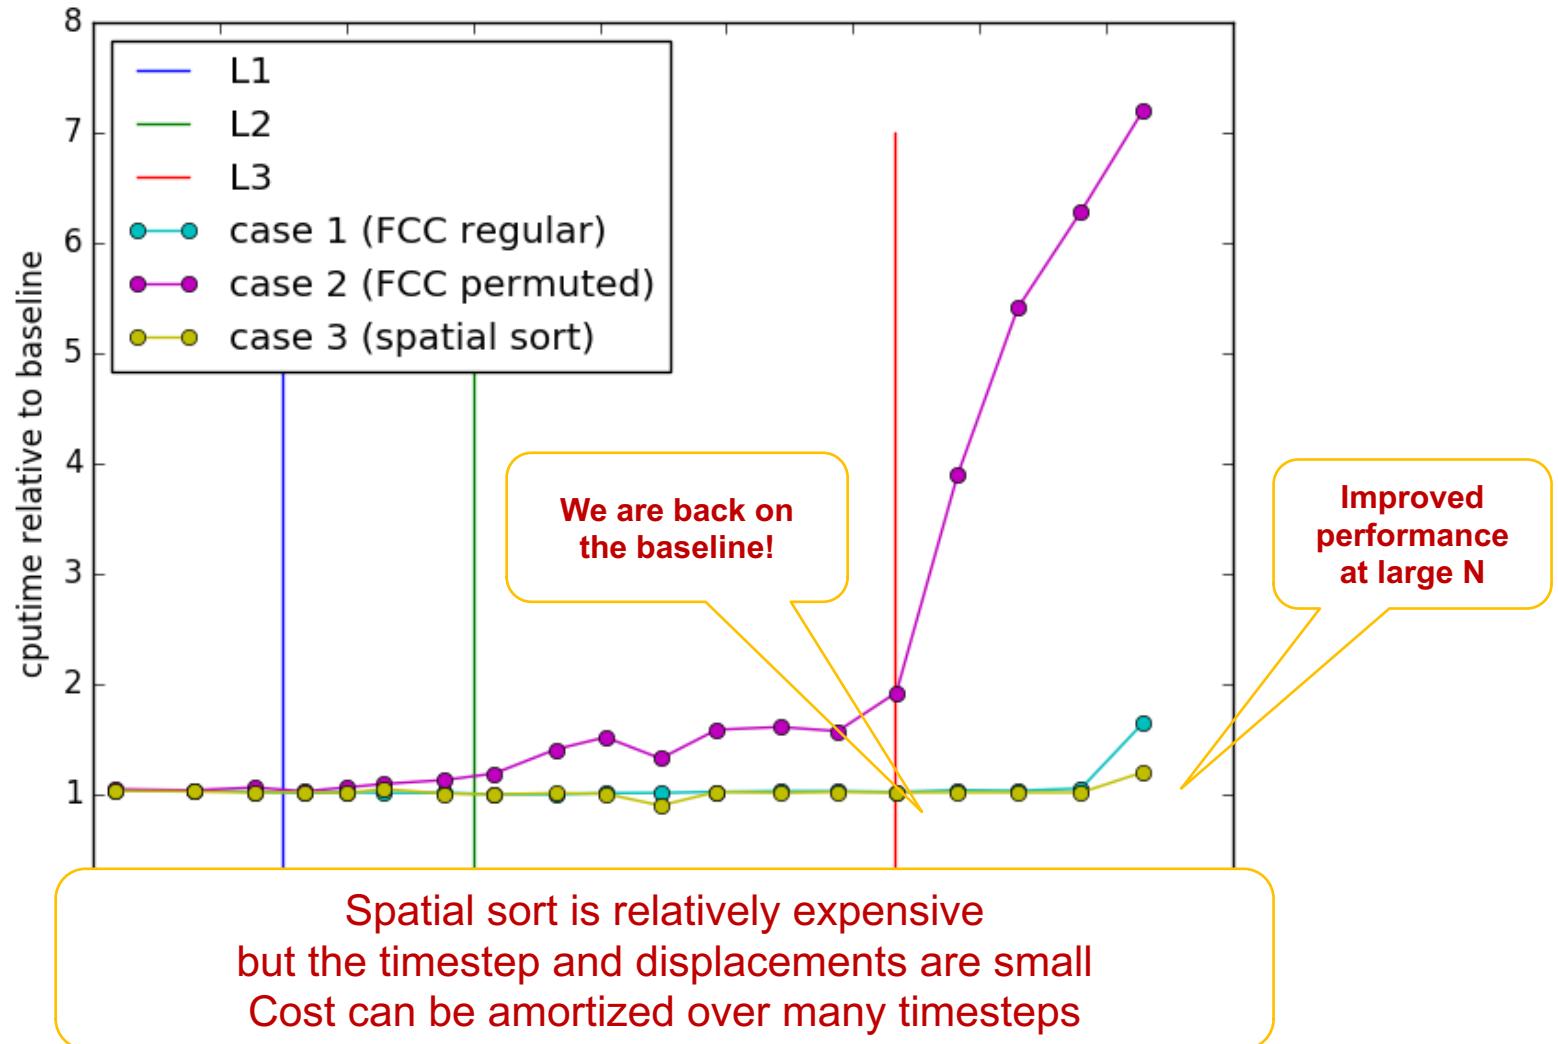
- Put atoms on FCC lattice
- 4 atoms per unit cell
- Closest neighbor distance = $LJ r_{min}$
- $r_{cutoff} = 3r_{min}$











Fixing the data access pattern

- 1. Sort atom property arrays ($rx, ry, rz, vx, vy, vz, \dots$) based on the Hilbert index h of the cell of the atoms (**spatial sort**).
Atoms which are close in space (and hence will interact) will be close in memory (and hence will be in the cache with high probability)
- 2. Build a table containing the index of the first atom in each cell, and the number of atoms in the cell (Hilbert list)
- 3. Build Verlet list from the Hilbert list (discard the latter)
- 4. Compute the interactions by looping over the Verlet list and **measure the performance** (e.g. interactions/s)
- 5. Integrate forces, updating velocities and positions and time
- 6. If **performance degrades**
 - jump back to step 1.
 - else
 - continue at step 4.

Performance analysis

- atoms : $2,13 \cdot 10^6$
- pairs : $162 \cdot 10^6$
- Ratio : 76
- Pairs computed per second : $88.6 \cdot 10^6$

- B/atom: 376,5
- Bandwidth = 9,5 GB/s
- maximum atoms per second: 27.000.000
- actual atoms per second: 1.160.000
- ratio: 0,04

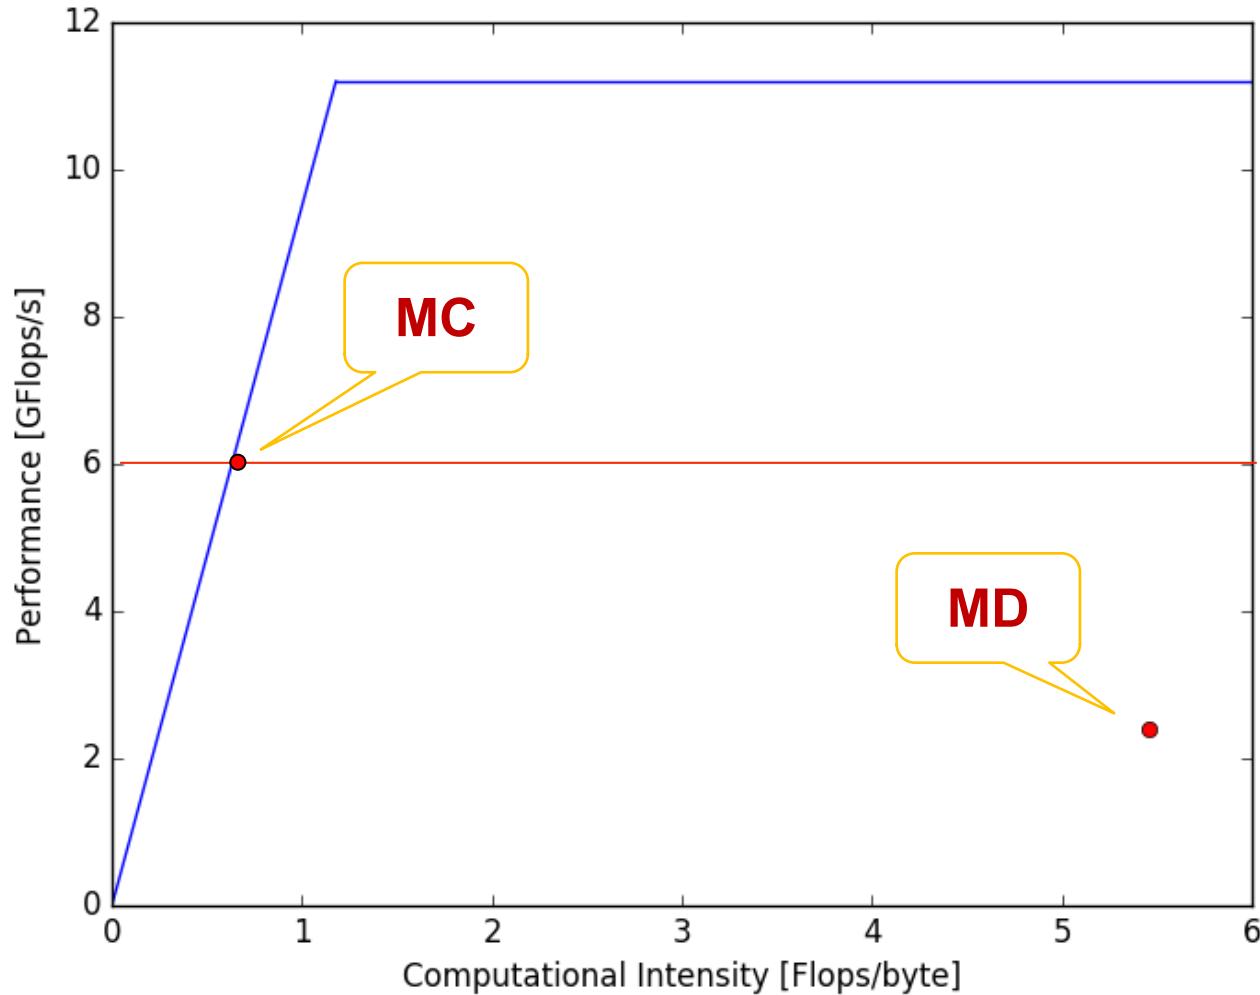
Not memory bound

- flops_per_pair : 27
- flops_per_atom : 2055,5

Lot of flops per atom!
- gflops_per_second: 2,39
- peak performance : 11,2
- ratio : 0,21

Not compute bound either!

Roofline MD setting



What do the intel tools tell? Advisor

Where should I add vectorization and/or threading parallelism?

Elapsed time: 55.75s Vectorized Not Vectorized FILTER: All Module All Sources Loops All Threads

Function Call Sites and Loops

Line	Source	Vector Issues	Self Time	Total Time	Type	Why No Vectorization?
317						
318	call cpu_time(cpustart)					
319	do iter=1,n_iter					
320	j=1					
321	do ia=1,n_atoms					
322	ia_pairs = verlet_linear(j)					
323	do k = j+1,j+ia_pairs					
	[loop in compute_interactions_verlet_linear at md.f90:323]	2 Assumed dependency present	9.390s	9.390s	Scalar	vector dependence prevents vectorization
			5.034s	5.034s	Scalar	

File: md.f90:328 compute_interactions_verlet_linear

Line	Source	Total Time	%	Loop Time	%	Traits
317						
318	call cpu_time(cpustart)					
319	do iter=1,n_iter					
320	j=1					
321	do ia=1,n_atoms					
322	ia_pairs = verlet_linear(j)					
323	do k = j+1,j+ia_pairs					
	[loop in compute_interactions_verlet_linear at md.f90:323]	0.010s	(17.900s	17.900s	Scalar loop. Not vectorized: vector dependence prevents vectorization No loop transformations applied
324	ja = verlet_linear(k)+1 ! +1 since fortran starts counting from 1 !					
325	dx = rx(ja)-rx(ia)					
326	dy = ry(ja)-ry(ia)					
327	dz = rz(ja)-rz(ia)					
328	aij = lj_force_factor2(dx**2 + dy**2 + dz**2)	0.280s	(9.729s	17.850s	17.850s
329	! update particle ia acceleration	0.330s	(
330	ax(ia) = ax(ia) + aij*dx	0.320s	(
331	ay(ia) = ay(ia) + aij*dy	0.527s	(
332	az(ia) = az(ia) + aij*dz	9.729s	17.850s			
333	! update particle ja acceleration	0.569s	(
334	ax(ja) = ax(ja) - aij*dx	2.758s	(
335	ay(ja) = ay(ja) - aij*dy	1.360s	(
336	az(ja) = az(ja) - aij*dz	0.190s	(
		0.688s	(

Selected (Total Time): 0s

What do the intel tools tell?

```
j=1
do ia=1,n_atoms
    ia_pairs = verlet_linear(j) ! Size of the Verlet list of atom ia
    do k = j+1,j+ia_pairs
        ja = verlet_linear(k)+1 ! +1 since Fortran starts counting from 1 !
        dx = rx(ja)-rx(ia)
        dy = ry(ja)-ry(ia)
        dz = rz(ja)-rz(ia)
        aij = lj_force_factor2( dx**2 + dy**2 + dz**2 )
        ! update particle ia acceleration
        ax(ia) = ax(ia) + aij*dx
        ay(ia) = ay(ia) + aij*dy
        az(ia) = az(ia) + aij*dz
        ! update particle ja acceleration
        ax(ja) = ax(ja) - aij*dx
        ay(ja) = ay(ja) - aij*dy
        az(ja) = az(ja) - aij*dz
    enddo
    j = j + 1 + ia_pairs
enddo
```

Ignore assumed dependencies and vectorize the loop

- SIMD vectorization means that you update $ax(aj)$ for 4 successive ja values (also $ay(aj)$ and $az(aj)$)
- The compiler cannot know that the ja are different
- Assumed dependency
- We know that the ja are different by construction of the Verlet list
- We must tell the compiler to ignore assumed dependencies

What do the intel tools tell?

- “inserts present” = hint for gather/scatter
- Filling a vector register element per element (in the case of non-contiguous elements)
- AVX (highest SIMD extension available on Hopper) has no built-in support for gather/scatter
- AVX2 has

```

j=1
do ia=1,n_atoms
    ia_pairs = verlet_linear(j) ! Size of the Verlet list of atom ia
    !DIR$ SIMD
    do k = j+1,j+ia_pairs
        ja = verlet_linear(k)+1 ! +1 since Fortran starts counting from 1 !
        dx = rx(ja) - rx(ia)
        dy = ry(ja) - ry(ia)
        dz = rz(ja) - rz(ia)
        aij = lj_force_factor2( dx**2 + dy**2 + dz**2 )
        ! update particle ia acceleration
        ax(ia) = ax(ia) + aij*dx
        ay(ia) = ay(ia) + aij*dy
        az(ia) = az(ia) + aij*dz
        ! update particle ja acceleration
        ax(ja) = ax(ja) - aij*dx
        ay(ja) = ay(ja) - aij*dy
        az(ja) = az(ja) - aij*dz
    enddo
    j = j + 1 + ia_pairs
enddo

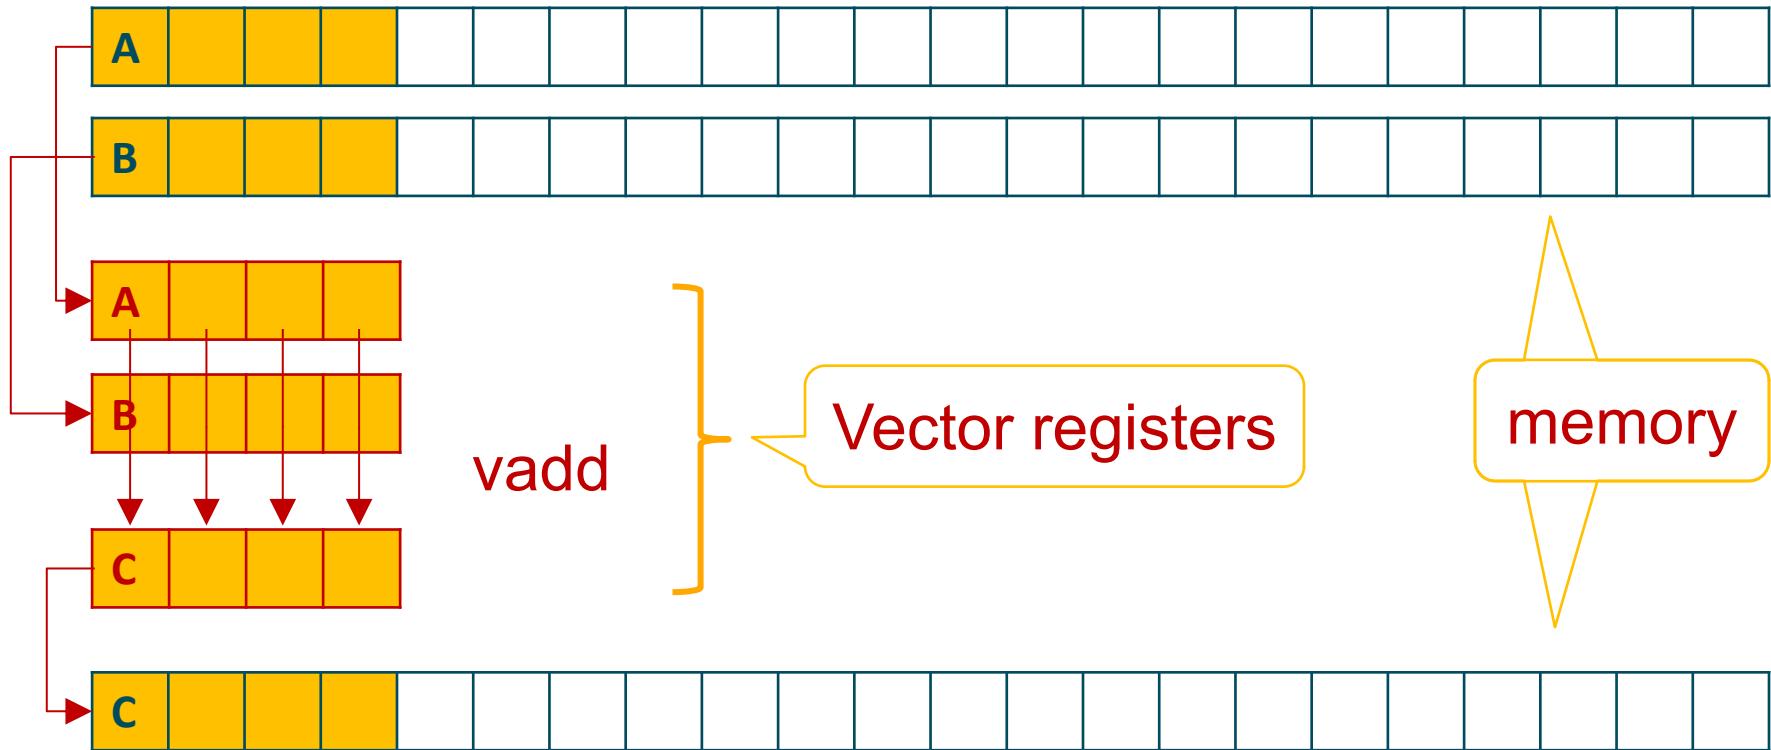
```

Gather operation moving $rx(ja)$, $ry(ja)$, $rz(ja)$ for 4 successive ja values into the vector registers

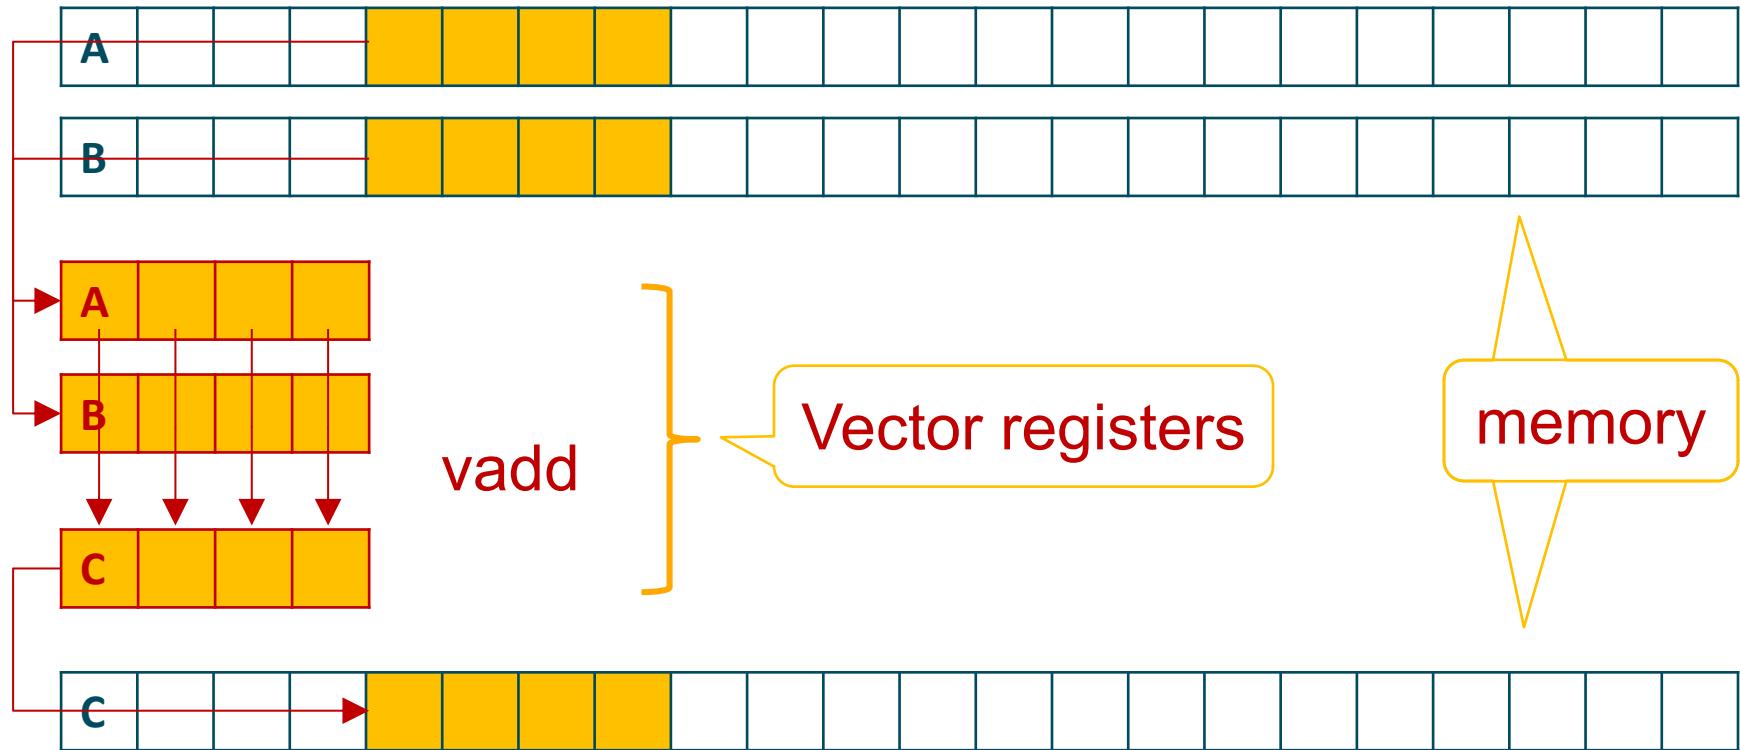
Scatter operation moving $ax(ja)$, $ay(ja)$, $az(ja)$ for 4 successive ja values out of the vector registers.

- Gather/scatter is also present in baseline!
- This is responsible for the extra speed loss factor of 1.66 relative to the MC setting:
- MC: **$450 \cdot 10^6$ interactions/s**
- $\div 3$ (3 times more data)
- **$150 \cdot 10^6$ interactions/s**
- $\div 1.66$ (gather/scatter)
- MD: **$90 \cdot 10^6$ interactions/s**

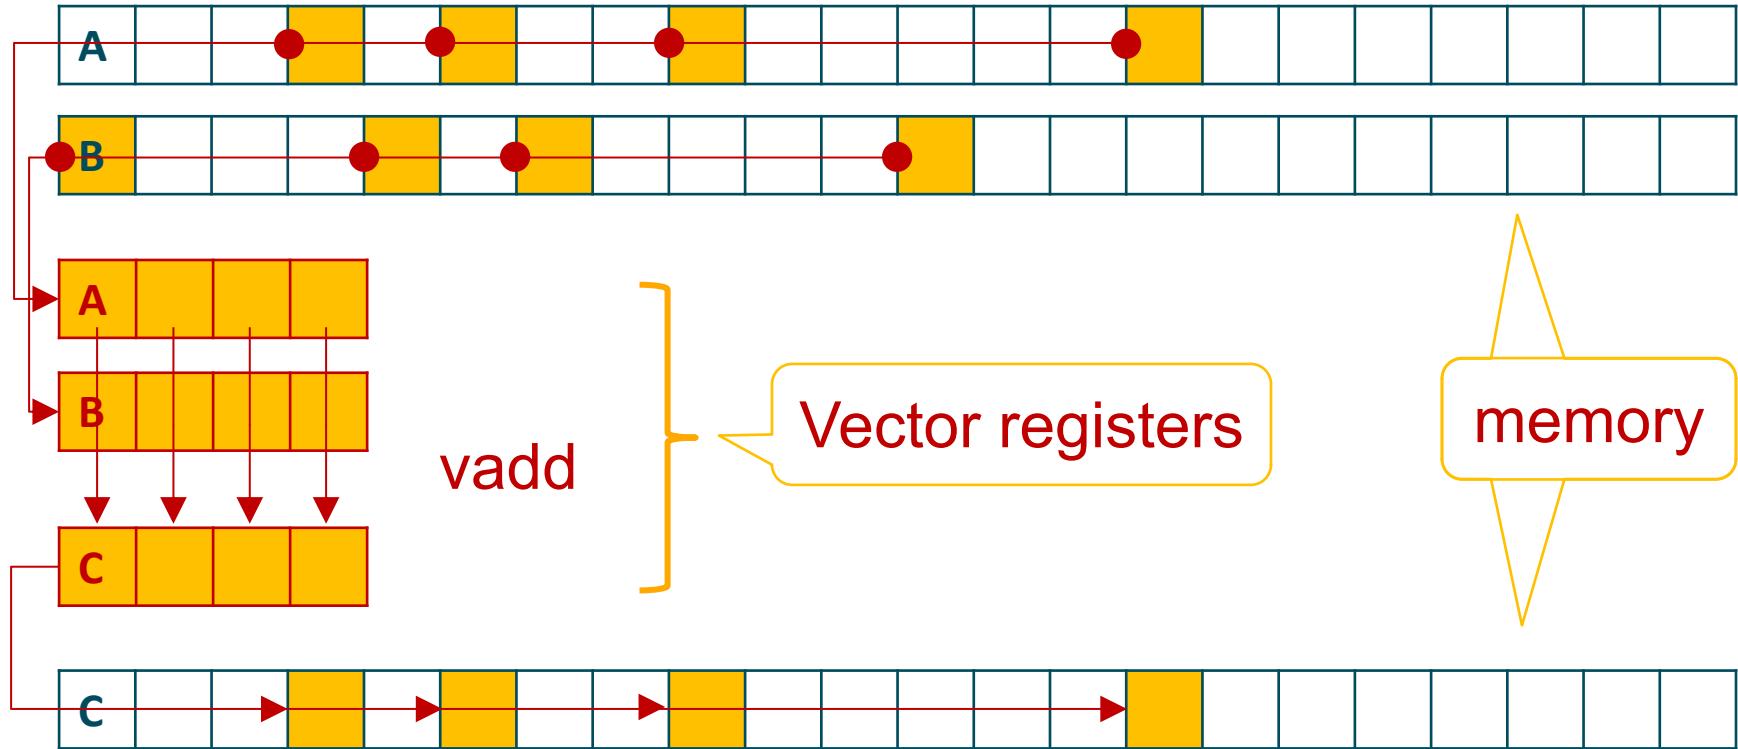
- $C = A+B$
- No gather/scatter, 4 successive items moved as a block into/out of vector registers



- $C = A+B$
- No gather/scatter, 4 successive items moved as a block into/out of vector registers



- $C = A+B$
- gather/scatter, items moved one by one into/out of vector registers



What do the intel tools tell?

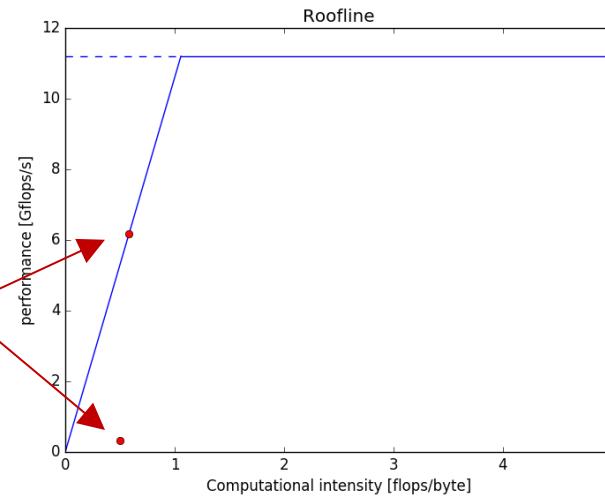
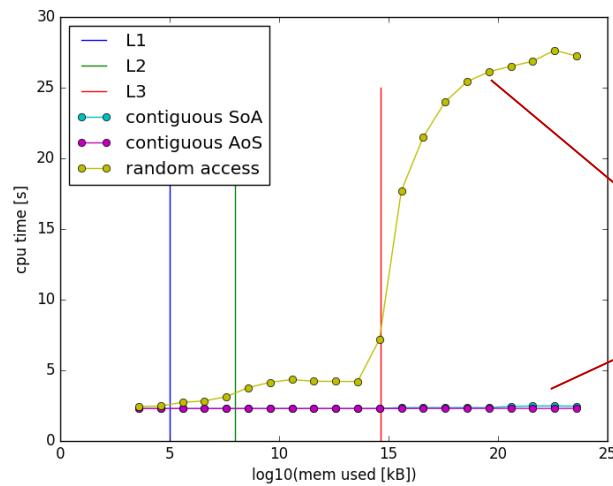
- LLC misses 0.008 %
 - LJ_force_factor2 0.59 cpi
 - Compute_interactions 0.77 cpi
-
- cpi at best 0.25
 - cpi ≤ 1 considered acceptable in HPC
-
- There is still room for improvement
 - No more low hanging fruit, though

Conclusions

1. Data access pattern is crucial to performance

Conclusions

- Monte Carlo case (3N reads, no writes)
- Contiguous data access:
 - $450 \cdot 10^6$ interactions/s
 - Bandwidth saturation (machine limit)
- Random access:
 - Performance drops by factor 15 (cache misses, gather/scatter)

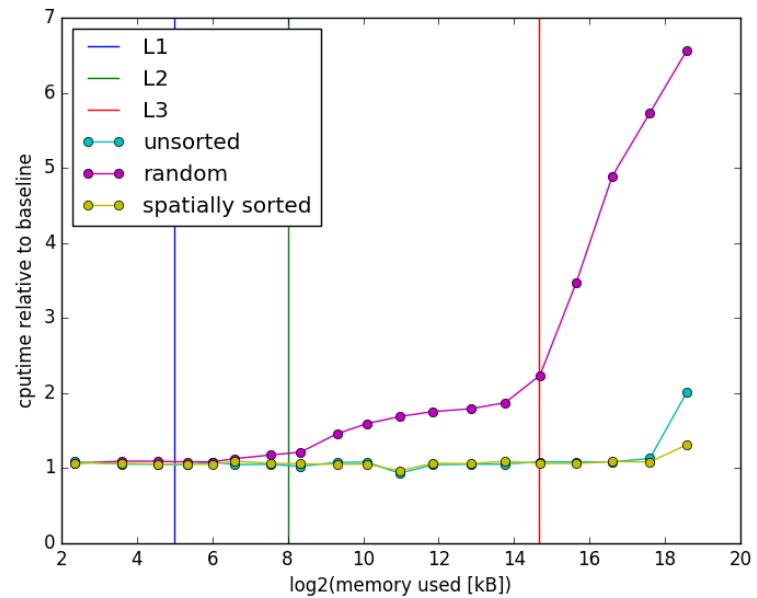


Conclusions

- Bandwidth saturation means that the CPU is waiting for the data to arrive
 - Try to do more computations with the data that is available, e.g.:
 - Additional computations
 - A more complex model
 - Program will not run faster but will do more work in the same time

Conclusions

- Molecular dynamics case ($6N$ reads, $3N$ writes + VL)
- Contiguous data access:
 - $90 \cdot 10^6$ interactions/s
 - No machine limits hit
 - gather/scatter
- Random access:
 - Performance drops by factor 7
- Spatial sort fixes the problem
- Fixing the data access pattern is not always easy
- Involves usually some form of sorting the data



1. Data access pattern is crucial to performance
2. **Spatial sort using space filling curves is useful technique for fixing data access patterns**

Conclusions

1. Data access pattern is crucial to performance
2. **Spatial sort using space filling curves is useful technique for fixing data access patterns**

Conclusions

1. Data access pattern is crucial to performance
2. Spatial sort using space filling curves
- 3. Intel tools (Advisor, VTune) provide useful clues to optimizing your code**
 - Hot spots
 - The nature of hot spots (data access, expensive instructions, ...)
 - Issues with vectorization

Conclusions

1. Data access pattern is crucial to performance
 2. Spatial sort using space filling curves is useful
 3. Intel tools (Advisor, VTune) provide useful clues
- 4. Fortran has many advantages for programming number-crunching routines**
- F2py for producing python modules

Python/Numpy

- verify code correctness
- generate FCC lattice
- generate arrays filled with random numbers
- generate permutations
- zero accelerations between time steps
- **build Verlet list**
- spatial sort of atom property arrays
- define coarse computational strategy and data structures
- control and initialize the experiments
- plot results

code **80%**

cputime 5%

Fortran

- Lennard-Jones potential
- Lennard-Jones forces
- iterate over Verlet list and compute interactions
- iterate over array and compute interactions (baselines)

Speedup of 1200x !

C++

- compute hilbert indices
- **build Verlet list**

code 10%

cputime **90%**

10%

cputime 5%

Conclusions

1. Data access pattern is crucial to performance
 2. Spatial sort using space filling curves is useful
 3. Intel tools (Advisor, VTune) provide useful clues
 4. Fortran has many advantages
- 5. Use a simple relevant baseline that you understand to judge the performance of your code and direct your efforts**

Conclusions

1. Data access pattern is crucial to performance
2. Spatial sort using space filling curves is useful
3. Intel tools (Advisor, VTune) provide useful clues
4. Fortran has many advantages
5. Use a simple relevant baseline

6. Be aware of the machine limits

- **Bandwidth**
- **Peak performance**
- **Roofline model**

Conclusions

1. Data access pattern is crucial to performance
2. Spatial sort using space filling curves is useful
3. Intel tools (Advisor, VTune) provide useful clues
4. Fortran has many advantages
5. Use a simple relevant baseline
6. Be aware of the machine limits

Thank you

You are always welcome to discuss your
(computational) problems